

Introduction

The garment industry refers to the production and distribution of clothing and related accessories, which also includes the textile sector that creates fibers and fabrics. As a significant part of the global manufacturing landscape, the garment industry employs millions of people and contributes trillions of dollars to the economy. Given its scale, maximizing workforce productivity is essential for the industry's success. This project focuses on estimating the productivity score for each team in a garment factory by analyzing a dataset related to worker productivity using an Artificial Neural Network (ANN) model in deep learning. ANN is particularly effective at identifying complex patterns within data, allowing us to pinpoint key factors that influence productivity scores. Through systematic data collection, preprocessing, model development, and evaluation, we aim to derive valuable insights that can improve efficiency and assist in strategic decision-making within the garment industry.

```
# Install keras tuner
```

```
!pip install keras-tuner
```

```
Requirement already satisfied: keras-tuner in  
/usr/local/lib/python3.11/dist-packages (1.4.7)  
Requirement already satisfied: keras in  
/usr/local/lib/python3.11/dist-packages (from keras-tuner) (3.8.0)  
Requirement already satisfied: packaging in  
/usr/local/lib/python3.11/dist-packages (from keras-tuner) (24.2)  
Requirement already satisfied: requests in  
/usr/local/lib/python3.11/dist-packages (from keras-tuner) (2.32.3)  
Requirement already satisfied: kt-legacy in  
/usr/local/lib/python3.11/dist-packages (from keras-tuner) (1.0.5)  
Requirement already satisfied: absl-py in  
/usr/local/lib/python3.11/dist-packages (from keras->keras-tuner)  
(1.4.0)  
Requirement already satisfied: numpy in  
/usr/local/lib/python3.11/dist-packages (from keras->keras-tuner)  
(2.0.2)  
Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-  
packages (from keras->keras-tuner) (13.9.4)  
Requirement already satisfied: namex in  
/usr/local/lib/python3.11/dist-packages (from keras->keras-tuner)  
(0.0.8)  
Requirement already satisfied: h5py in /usr/local/lib/python3.11/dist-  
packages (from keras->keras-tuner) (3.13.0)  
Requirement already satisfied: optree in  
/usr/local/lib/python3.11/dist-packages (from keras->keras-tuner)  
(0.15.0)  
Requirement already satisfied: ml-dtypes in  
/usr/local/lib/python3.11/dist-packages (from keras->keras-tuner)  
(0.4.1)  
Requirement already satisfied: charset-normalizer<4,>=2 in
```

```
/usr/local/lib/python3.11/dist-packages (from requests->keras-tuner)
(3.4.1)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.11/dist-packages (from requests->keras-tuner)
(3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.11/dist-packages (from requests->keras-tuner)
(2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.11/dist-packages (from requests->keras-tuner)
(2025.1.31)
Requirement already satisfied: typing-extensions>=4.5.0 in
/usr/local/lib/python3.11/dist-packages (from optree->keras->keras-
tuner) (4.13.1)
Requirement already satisfied: markdown-it-py>=2.2.0 in
/usr/local/lib/python3.11/dist-packages (from rich->keras->keras-
tuner) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in
/usr/local/lib/python3.11/dist-packages (from rich->keras->keras-
tuner) (2.18.0)
Requirement already satisfied: mdurl~=0.1 in
/usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0-
>rich->keras->keras-tuner) (0.1.2)
```

```
# Import Library
```

```
import keras_tuner as kt
```

```
import pandas as pd
```

```
import numpy as np
```

```
import random
```

```
import tensorflow as tf
```

```
from tensorflow.keras.layers import Dense, Dropout, Input
```

```
from tensorflow.keras.models import Sequential, Model
```

```
from tensorflow.keras.utils import plot_model
```

```
from tensorflow.keras.regularizers import l2
```

```
from tensorflow.keras.optimizers import Adam
```

```
from tensorflow.keras.callbacks import EarlyStopping
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import
```

```
RobustScaler, StandardScaler, MinMaxScaler, LabelEncoder, OneHotEncoder
```

```
from sklearn.impute import SimpleImputer
```

```
from sklearn.metrics import r2_score, mean_absolute_error,
mean_squared_error
```

```
from warnings import filterwarnings
```

```
filterwarnings('ignore')
```

```
SEED_VALUE=1234
```

```
random.seed(SEED_VALUE)
```

```
np.random.seed(SEED_VALUE)
```

```
# Load Dataset
```

```
df = pd.read_parquet("dataset_1B.parquet")
```

```
df.head()
```

```
{"summary":{"name": "df", "rows": 1197, "fields":  
[\n  {\n    "column": "date", "properties": {\n      "dtype": "category", "num_unique_values": 118, \n      "samples": [\n        "2/3/2015", "2016-02-23", \n        "1/4/2015"\n      ], \n      "semantic_type": "", \n      "description": ""\n    }, \n    "column": "quarter", "properties": {\n      "dtype": "category", \n      "num_unique_values": 5, \n      "samples": [\n        "Quarter2", "Quarter5", \n        "Quarter3"\n      ], \n      "semantic_type": "", \n      "description": ""\n    }, \n    "column": "day", "properties": {\n      "dtype": "category", \n      "num_unique_values": 6, \n      "samples": [\n        "Thursday", "Saturday", "Wednesday"\n      ], \n      "semantic_type": "", \n      "description": ""\n    }, \n    "column": "Team Code", "properties": {\n      "dtype": "number", "std": 3, \n      "min": 1, "max": 12, \n      "num_unique_values": 12, \n      "samples": [\n        5, \n        10, \n        8\n      ], \n      "semantic_type": "", \n      "description": ""\n    }, \n    "column": "smv", "properties": {\n      "dtype": "number", \n      "std": 10.943219199514333, \n      "min": 2.9, \n      "max": 54.56, \n      "num_unique_values": 70, \n      "samples": [\n        14.61, \n        26.16, \n        30.1\n      ], \n      "semantic_type": "", \n      "description": ""\n    }, \n    "column": "wip", "properties": {\n      "dtype": "number", \n      "std": 1837.4550011056342, \n      "min": 7.0, \n      "max": 23122.0, \n      "num_unique_values": 548, \n      "samples": [\n        1287.0, \n        970.0, \n        958.0\n      ], \n      "semantic_type": "", \n      "description": ""\n    }, \n    "column": "over_time", "properties": {\n      "dtype": "number", \n      "std": 3348, \n      "min": 0, \n      "max": 25920, \n      "num_unique_values": 143, \n      "samples": [\n        5820, \n        6780, \n        7470\n      ], \n      "semantic_type": "", \n      "description": ""\n    }, \n    "column": "incentive", "properties": {\n      "dtype": "number", \n      "std": 160, \n      "min": 0,
```

```

{"max": 3600, "num_unique_values": 48, "samples": [55, 65, 81], "semantic_type": "", "description": "", "column": "idle_time", "properties": {"dtype": "number", "std": 12.709756518546547, "min": 0.0, "max": 300.0, "num_unique_values": 12, "samples": [4.0, 3.5, 0.0], "semantic_type": "", "description": "", "column": "idle_men", "properties": {"dtype": "number", "std": 3, "min": 0, "max": 45, "num_unique_values": 10, "samples": [25, 10, 30], "semantic_type": "", "description": "", "column": "no_of_style_change", "properties": {"dtype": "number", "std": 0, "min": 0, "max": 2, "num_unique_values": 3, "samples": [0, 1, 2], "semantic_type": "", "description": "", "column": "no_of_workers", "properties": {"dtype": "number", "std": 22.61704313753719, "min": -57.0, "max": 89.0, "num_unique_values": 66, "samples": [50.0, 46.0, 59.0], "semantic_type": "", "description": "", "column": "productivity_score", "properties": {"dtype": "number", "std": 18.154945385667887, "min": -100.0, "max": 112.044, "num_unique_values": 803, "samples": [97.462, 97.008, 80.031], "semantic_type": "", "description": ""}]}, "type": "dataframe", "variable_name": "df"}

```

Structure

```
# Shape of the dataset
```

```
df.shape
```

```
(1197, 13)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1197 entries, 0 to 1196
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----

0	date	1197	non-null	object
1	quarter	1197	non-null	object
2	day	1197	non-null	object
3	Team Code	1197	non-null	int64
4	smv	1197	non-null	float64
5	wip	691	non-null	float64
6	over_time	1197	non-null	int64
7	incentive	1197	non-null	int64
8	idle_time	1197	non-null	float64
9	idle_men	1197	non-null	int64
10	no_of_style_change	1197	non-null	int64
11	no_of_workers	1197	non-null	float64
12	productivity_score	1197	non-null	float64

dtypes: float64(5), int64(5), object(3)
memory usage: 121.7+ KB

The dataset consists of 1197 entries (rows) and 13 columns (features), as indicated by its shape of (1197, 13).

Based on the information of the dataset, It includes three categorical columns—date, quarter, and day, along with 10 numerical columns such as Team Code, smv (Standard Minute Value), wip (Work In Progress), over_time, incentive, idle_time, idle_men, no_of_style_change, no_of_workers, and productivity_score. Further analysis is needed to determine the appropriate handling of these variables.

Below are the detailed features in the dataset:

- date: The date of assessment
- day: The day of the Week
- quarter: The quarter of the year when the data was recorded (e.g., Quarter1, Quarter2)
- Team Code: A unique identifier for the team.
- smv: The Standard Minute Value, which indicates the time allocated for specific task.
- wip: Work In Progress, reflecting number of unfinished products
- over_time: The amount of overtime worked, measured in minutes.
- incentive: The financial incentive provided to the workers, measured in USD.
- idle_time: The total time that workers were idle, also measured in minutes.
- idle_men: The number of workers who were idle / not engaged in work.
- no_of_style_change: The total number of style changes that occurred.
- no_of_workers: The total number of workers in the team.
- productivity_score: The team's productivity score, expressed as a percentage.

```
# Display all columns in the dataset
df.columns
```

```
Index(['date', 'quarter', 'day', 'Team Code', 'smv', 'wip',
      'over_time',
      'incentive', 'idle_time', 'idle_men', 'no_of_style_change',
```



```
89.0,\n          54.56\n          ],\n          \"semantic_type\": \"\", \n          \"description\": \"\" \n          } \n          } \n          ] \n          }\", \"type\": \"dataframe\"}
```

The summary statistics of the dataset reveal important insights into various numerical features. The Team Code has a mean of 6.42 with a limited range, indicating a small number of teams. Meanwhile, the Incentive varies widely, with a mean of 38.21 USD, although some teams receive none. Additionally, the average idle time and number of idle men are low, suggesting minimal inefficiencies. However, there are negative values in the "no_of_workers" column (-57) and the "productivity_score" column (-100), which need to be addressed during preprocessing to clean the data. Overall, these statistics guide further investigations and analyses aimed at optimizing productivity in the garment factory.

Data Cleaning

```
# Check duplicate data
df.duplicated().sum()

np.int64(0)
```

By checking duplicate data in the dataset, it was confirmed that there are no duplicates, ensuring that it is free from redundancy and ready for further analysis.

```
# Check missing value
df.isna().sum()

date                0
quarter             0
day                 0
Team Code           0
smv                 0
wip                 506
over_time           0
incentive           0
idle_time           0
idle_men            0
no_of_style_change  0
no_of_workers       0
productivity_score  0
dtype: int64
```

After checking the missing values for each column, it was found that "wip" (work in progress) column has 506 missing values. Meanwhile, other columns contain no missing data.

This highlights the importance of addressing missing values in "wip" column, that require us to decide whether impute or use other handling methods when preprocessing.

```
# Rename column
```

```
df = df.rename(columns={'wip': 'work_in_progress', 'smv':  
'std_minute_value'})  
df.head()
```

```
{"summary":{"\n  \"name\": \"df\",\n  \"rows\": 1197,\n  \"fields\":  
[\n    {\n      \"column\": \"date\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 118,\n        \"samples\": [\n          \"2/3/2015\",\n          \"2016-02-23\",\n          \"1/4/2015\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\",\n        \"quarter\": \"\",\n        \"properties\": {\n          \"dtype\": \"category\",\n          \"num_unique_values\": 5,\n          \"samples\": [\n            \"Quarter2\",\n            \"Quarter5\",\n            \"Quarter3\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\",\n          \"day\": \"\",\n          \"properties\": {\n            \"dtype\": \"category\",\n            \"num_unique_values\": 6,\n            \"samples\": [\n              \"Thursday\",\n              \"Saturday\",\n              \"Wednesday\"\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          }\n        },\n        {\n          \"column\": \"Team Code\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 3,\n            \"min\": 1,\n            \"max\": 12,\n            \"num_unique_values\": 12,\n            \"samples\": [\n              5,\n              10,\n              8\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\",\n            \"std_minute_value\": \"\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 10.943219199514333,\n              \"min\": 2.9,\n              \"max\": 54.56,\n              \"num_unique_values\": 70,\n              \"samples\": [\n                14.61,\n                26.16,\n                30.1\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            }\n          },\n          {\n            \"column\": \"work_in_progress\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 1837.4550011056342,\n              \"min\": 7.0,\n              \"max\": 23122.0,\n              \"num_unique_values\": 548,\n              \"samples\": [\n                1287.0,\n                970.0,\n                958.0\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            }\n          },\n          {\n            \"column\": \"over_time\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 3348,\n              \"min\": 0,\n              \"max\": 25920,\n              \"num_unique_values\": 143,\n              \"samples\": [\n                5820,\n                6780,\n                7470\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            }\n          },\n          {\n            \"column\": \"incentive\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 160,\n              \"min\": 0,\n              \"max\": 3600,\n              \"num_unique_values\": 48,\n              \"samples\": [\n                55,\n                65,\n                81\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            }\n          },\n          {\n            \"column\": \"idle_time\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 160,\n              \"min\": 0,\n              \"max\": 3600,\n              \"num_unique_values\": 48,\n              \"samples\": [\n                55,\n                65,\n                81\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            }\n          }\n        ]\n      }\n    }\n  ]\n}
```



```

{"number": 0.0, "std": 12.709756518546547, "min": 0.0, "max": 300.0, "num_unique_values": 12, "samples": [4.0, 3.5, 0.0], "semantic_type": "", "description": ""}, {"column": "idle_men", "properties": {"dtype": "number", "std": 3, "min": 0, "max": 45, "num_unique_values": 10, "samples": [25, 10, 30], "semantic_type": "", "description": ""}, {"column": "no_of_style_change", "properties": {"dtype": "number", "std": 0, "min": 0, "max": 2, "num_unique_values": 3, "samples": [0, 1, 2], "semantic_type": "", "description": ""}, {"column": "no_of_workers", "properties": {"dtype": "number", "std": 22.61704313753719, "min": -57.0, "max": 89.0, "num_unique_values": 66, "samples": [50.0, 46.0, 59.0], "semantic_type": "", "description": ""}, {"column": "productivity_score", "properties": {"dtype": "number", "std": 18.154945385667887, "min": -100.0, "max": 112.044, "num_unique_values": 803, "samples": [97.462, 97.008, 80.031], "semantic_type": "", "description": ""}]
n}, {"type": "dataframe", "variable_name": "df"}

```

I decided to rename the "wip" column to "work_in_progress" and the "smv" column to "std_minute_value" to make them easier to understand.

```

# Convert "date" column into datetime format
df['date'] = pd.to_datetime(df['date'], format = 'mixed')

df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1197 entries, 0 to 1196
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  1197 non-null  datetime64[ns]
1   quarter               1197 non-null  object
2   day                   1197 non-null  object
3   Team Code             1197 non-null  int64
4   std_minute_value      1197 non-null  float64
5   work_in_progress      691 non-null   float64
6   over_time             1197 non-null  int64

```

```

7    incentive          1197 non-null    int64
8    idle_time          1197 non-null    float64
9    idle_men           1197 non-null    int64
10   no_of_style_change 1197 non-null    int64
11   no_of_workers       1197 non-null    float64
12   productivity_score  1197 non-null    float64
dtypes: datetime64[ns](1), float64(5), int64(5), object(2)
memory usage: 121.7+ KB

# Split "date" column into 3 new column and drop "date" column
df['Date'] = df['date'].dt.day
df['Month'] = df['date'].dt.month
df['Year'] = df['date'].dt.year

df = df.drop(columns = 'date')

```

Next, I decided to convert the "date" column from an object to a datetime format to ensure the data type is correct. Additionally, I split the "date" column into 3 separate columns which are "Day", "Month" and "Year" to maintain clarity regarding the "quarter" column are related to months in a year. This approach helps to avoid any confusion in the analysis.

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1197 entries, 0 to 1196
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype
---  -
0   quarter               1197 non-null   object
1   day                   1197 non-null   object
2   Team Code             1197 non-null   int64
3   std_minute_value      1197 non-null   float64
4   work_in_progress      691 non-null    float64
5   over_time             1197 non-null   int64
6   incentive             1197 non-null   int64
7   idle_time             1197 non-null   float64
8   idle_men              1197 non-null   int64
9   no_of_style_change    1197 non-null   int64
10  no_of_workers          1197 non-null   float64
11  productivity_score     1197 non-null   float64
12  Date                   1197 non-null   int32
13  Month                  1197 non-null   int32
14  Year                   1197 non-null   int32
dtypes: float64(5), int32(3), int64(5), object(2)
memory usage: 126.4+ KB

```

Now, the "date" column has been dropped, and the 3 newly created column are set in numerical format.

```
# Numerical Analysis
```

```
df.describe().T
```

```
{"summary":{"\n  \"name\": \"df\",\n  \"rows\": 13,\n  \"fields\": [\n    {\n      \"column\": \"count\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 140.33914964498294,\n        \"min\": 691.0,\n        \"max\": 1197.0,\n        \"num_unique_values\": 2,\n        \"samples\": [\n          691.0,\n          1197.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"mean\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1337.8633025679371,\n        \"min\": 0.15037593984962405,\n        \"max\": 4567.460317460317,\n        \"num_unique_values\": 13,\n        \"samples\": [\n          1.717627401837928,\n          73.36704010025063\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"std\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1014.5339489507038,\n        \"min\": 0.4278478565061976,\n        \"max\": 3348.823562883226,\n        \"num_unique_values\": 13,\n        \"samples\": [\n          0.7375261810563015,\n          18.154945385667887\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"min\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 563.0484819950437,\n        \"min\": -100.0,\n        \"max\": 2015.0,\n        \"num_unique_values\": 7,\n        \"samples\": [\n          1.0,\n          2.9\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"25%\"\n      },\n      \"column\": \"25%\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 664.7999150872388,\n        \"min\": 0.0,\n        \"max\": 2015.0,\n        \"num_unique_values\": 10,\n        \"samples\": [\n          1.0,\n          3.94\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"50%\"\n      },\n      \"column\": \"50%\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1186.1962432085065,\n        \"min\": 0.0,\n        \"max\": 3960.0,\n        \"num_unique_values\": 10,\n        \"samples\": [\n          2.0,\n          15.26\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"75%\"\n      },\n      \"column\": \"75%\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1950.2954183206284,\n        \"min\": 0.0,\n        \"max\": 6960.0,\n        \"num_unique_values\": 11,\n        \"samples\": [\n          0.0,\n          9.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"max\"\n      },\n      \"column\": \"max\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 9075.24221067502,\n        \"min\": 2.0,\n        \"max\": 25920.0,\n        \"num_unique_values\": 13,\n        \"samples\": [\n          3.0,\n          112.044\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    ]\n  }},\n  \"type\": \"dataframe\"}
```

```
df['no_of_workers'] = df['no_of_workers'].abs()  
df['productivity_score'] = df['productivity_score'].abs()
```

```
df.describe().T
```

```
{
  "summary": {
    "name": "df",
    "rows": 13,
    "fields": [
      {
        "column": "count",
        "properties": {
          "dtype": "number",
          "std": 140.33914964498294,
          "min": 691.0,
          "max": 1197.0,
          "num_unique_values": 2,
          "samples": [
            691.0,
            1197.0
          ],
          "semantic_type": "",
          "description": " ",
          "mean": 944.0,
          "properties": {
            "dtype": "number",
            "std": 1337.8478970353888,
            "min": 0.15037593984962405,
            "max": 4567.460317460317,
            "num_unique_values": 13,
            "samples": [
              1.717627401837928,
              73.53412447786133
            ],
            "semantic_type": "",
            "description": " ",
            "std": 1014.5701984729498,
            "min": 0.4278478565061976,
            "max": 3348.823562883226,
            "num_unique_values": 13,
            "samples": [
              0.7375261810563015,
              17.46529787377751
            ],
            "semantic_type": "",
            "description": " ",
            "min": 0.0,
            "max": 2015.0,
            "num_unique_values": 7,
            "samples": [
              1.0,
              2.9
            ],
            "semantic_type": "",
            "description": " ",
            "min": 0.0,
            "max": 2015.0,
            "num_unique_values": 10,
            "samples": [
              1.0,
              3.94
            ],
            "semantic_type": "",
            "description": " ",
            "min": 0.0,
            "max": 3960.0,
            "num_unique_values": 10,
            "samples": [
              2.0,
              15.26
            ],
            "semantic_type": "",
            "description": " ",
            "min": 0.0,
            "max": 6960.0,
            "num_unique_values": 11,
            "samples": [
              0.0,
              9.0
            ],
            "semantic_type": "",
            "description": " ",
            "min": 2.0,
            "max": 25920.0,
            "num_unique_values": 13,
            "samples": [
              3.0,
              112.044
            ],
            "semantic_type": "",
            "description": " "
          }
        }
      }
    ]
  },
  "type": "dataframe"
}
```

I decided to convert the negative values in the "no_of_workers" column and "productivity_score" column to ensure that all entries are positive, which maintains the integrity of the data and allows for accurate analysis in predicting productivity in the garment factory. This is crucial because the number of workers cannot be negative, and a productivity score should not be negative either, as it represents performance relative to a defined standard.

EDA (Exploratory Data Analysis)

```
# Missing values
df.isna().sum()

quarter          0
day              0
Team Code        0
std_minute_value 0
work_in_progress 506
over_time        0
incentive        0
idle_time        0
idle_men         0
no_of_style_change 0
no_of_workers    0
productivity_score 0
Date             0
Month            0
Year            0
dtype: int64

# Check skew value for the work_in_progress column
df.work_in_progress.skew()

np.float64(9.741786273952965)

# Fill missing value with median
df['work_in_progress'].fillna(df['work_in_progress'].median(),inplace=True)
df.work_in_progress.median()

1039.0
```

It was observed that there are 506 missing values in the "work_in_progress" column. I checked the skewness to understand the distribution of the data, which revealed a skewness of 9.74, indicating a highly right-skewed distribution. This suggests that there are a significant number of lower values, with a few high values pulling the mean upward. Consequently, to determine the best method for imputing the missing values, I decided to use the median, because it is a robust statistic that is less influenced by outliers, making it a more accurate representation of the central tendency in this case.

```
# Re-check missing values
```

```
df.isna().sum()
```

```
quarter      0
day           0
Team Code     0
std_minute_value  0
work_in_progress  0
over_time    0
incentive     0
idle_time     0
idle_men      0
no_of_style_change  0
no_of_workers  0
productivity_score  0
Date          0
Month         0
Year          0
dtype: int64
```

Now, the data has no missing values and is clean, making it ready for further analysis. This preparation allows us to confidently proceed with tasks such as modeling and predicting productivity scores within the garment dataset.

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1197 entries, 0 to 1196
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null	Count	Dtype
0	quarter	1197	non-null	object
1	day	1197	non-null	object
2	Team Code	1197	non-null	int64
3	std_minute_value	1197	non-null	float64
4	work_in_progress	1197	non-null	float64
5	over_time	1197	non-null	int64
6	incentive	1197	non-null	int64
7	idle_time	1197	non-null	float64
8	idle_men	1197	non-null	int64
9	no_of_style_change	1197	non-null	int64
10	no_of_workers	1197	non-null	float64
11	productivity_score	1197	non-null	float64
12	Date	1197	non-null	int32
13	Month	1197	non-null	int32
14	Year	1197	non-null	int32

```
dtypes: float64(5), int32(3), int64(5), object(2)
```

```
memory usage: 126.4+ KB
```

```
# Separating categorical and numerical features based on dataset info
cat_cols = ['quarter', 'day']
num_cols = ['std_minute_value', 'over_time', 'incentive', 'idle_time',
'idle_men', 'no_of_style_change', 'no_of_workers',
'productivity_score', 'Date', 'Month', 'Year']
```

```
# Value count for each columns
```

```
for i in df.columns:
    x = df[i].value_counts()
    print(f'{x}\n')
```

```
quarter
Quarter1    360
Quarter2    335
Quarter4    248
Quarter3    210
Quarter5     44
Name: count, dtype: int64
```

```
day
Wednesday    208
Sunday        203
Tuesday       201
Thursday      199
Monday        199
Saturday      187
Name: count, dtype: int64
```

```
Team Code
8      109
2      109
4      105
1      105
9      104
10     100
12      99
7       96
3       95
6       94
5       93
11      88
Name: count, dtype: int64
```

```
std_minute_value
3.94      192
2.90      108
22.52     103
30.10      79
4.15       76
...
```

```
38.09      1
48.18      1
30.40      1
50.89      1
20.20      1
Name: count, Length: 70, dtype: int64
```

```
work_in_progress
1039.0     511
1282.0       4
1144.0       3
1193.0       3
1069.0       3
...
817.0       1
1576.0       1
1262.0       1
953.0        1
1161.0       1
Name: count, Length: 548, dtype: int64
```

```
over_time
960       129
1440      111
6960       61
6840       48
1200       39
...
5700       1
1680       1
1700       1
4680       1
3120       1
Name: count, Length: 143, dtype: int64
```

```
incentive
0         604
50        113
63         61
45         54
30         52
23         38
38         29
60         28
40         27
75         24
113        21
88         19
34         17
56         14
```


26	9
55	7
81	7
100	7
65	6
69	6
70	6
960	5
35	5
44	4
94	4
90	3
49	2
27	2
46	2
119	2
24	2
98	1
29	1
54	1
37	1
21	1
138	1
33	1
53	1
93	1
62	1
32	1
1080	1
2880	1
3600	1
1440	1
1200	1
25	1

Name: count, dtype: int64

idle_time	
0.0	1179
3.5	3
2.0	2
8.0	2
4.0	2
4.5	2
5.0	2
90.0	1
270.0	1
150.0	1
300.0	1
6.5	1

```
Name: count, dtype: int64
```

```
idle_men
```

```
0      1179
```

```
10       3
```

```
15       3
```

```
30       3
```

```
20       3
```

```
35       2
```

```
37       1
```

```
45       1
```

```
25       1
```

```
40       1
```

```
Name: count, dtype: int64
```

```
no_of_style_change
```

```
0      1050
```

```
1       114
```

```
2        33
```

```
Name: count, dtype: int64
```

```
no_of_workers
```

```
8.0      262
```

```
58.0     114
```

```
57.0     109
```

```
59.0      75
```

```
10.0      60
```

```
...
```

```
26.0       1
```

```
47.0       1
```

```
48.0       1
```

```
24.0       1
```

```
6.0        1
```

```
Name: count, Length: 61, dtype: int64
```

```
productivity_score
```

```
80.040     25
```

```
75.065     17
```

```
80.012     13
```

```
85.014     12
```

```
97.187     12
```

```
..
```

```
37.047      1
```

```
97.756      1
```

```
94.560      1
```

```
84.053      1
```

```
77.815      1
```

```
Name: count, Length: 802, dtype: int64
```

```
Date
```

```

10    65
7     64
8     62
3     61
4     59
1     58
11    54
5     52
12    42
24    42
28    42
22    42
25    42
17    41
2     41
9     41
26    40
18    40
19    39
15    38
14    38
6     28
29    28
31    24
13    22
27    21
23    19
21    19
16    18
20    15

```

Name: count, dtype: int64

Month

```

1     542
2     451
3     204

```

Name: count, dtype: int64

Year

```

2015    701
2016    496

```

Name: count, dtype: int64

I performed value counts for each column in the dataset to analyze the distribution of categorical and numerical features. In the "Quarter" column, Quarter 1 had the highest frequency with 360 entries, while the "Day" column showed Wednesday as the most common day with 208 entries. The "Team Code" values were fairly balanced, with Team 8 and Team 2 each having 109 entries. The "Standard Minute Value" ranged across 70 unique values, with the most frequent being 3.94. The "Work In Progress" column showed that 1039.0 was the most

common value, and the majority of entries in "Idle Time" and "Idle Men" reported zero. The "Incentive" column revealed that 604 entries received no incentive, while the "Productivity Score" had 803 unique values. Overall, this analysis provides valuable insights into the dataset, aiding in understanding trends and patterns related to productivity in the garment factory.

```
# Check unique value in categorical column
```

```
for i in cat_cols:
    print(f'{i}: {df[i].nunique()}')
    print(df[i].unique())
    print()
```

```
quarter: 5
```

```
['Quarter1' 'Quarter2' 'Quarter3' 'Quarter4' 'Quarter5']
```

```
day: 6
```

```
['Thursday' 'Saturday' 'Sunday' 'Monday' 'Tuesday' 'Wednesday']
```

I checked the unique values in the categorical columns, and it showed that there are 5 quarters. This indicates that Quarter 5 needs to be addressed since there should only be 4 quarters in total. Additionally, in the "day" column, there are only 6 days represented, with no entries for Friday.

```
# Filter the dataset for any records associated with that specific quarter
```

```
df[df['quarter'] == 'Quarter5']
```

```
{"summary": "{\n  \"name\": \"df[df['quarter'] == 'Quarter5']\",\n  \"rows\": 44,\n  \"fields\": [\n    {\n      \"column\": \"quarter\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 1,\n        \"samples\": [\n          \"Quarter5\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"day\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 2,\n        \"samples\": [\n          \"Saturday\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Team Code\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 3,\n        \"min\": 1,\n        \"max\": 12,\n        \"num_unique_values\": 12,\n        \"samples\": [\n          12\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"std_minute_value\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 10.94775058469517,\n        \"min\": 2.9,\n        \"max\": 50.89,\n        \"num_unique_values\": 14,\n        \"samples\": [\n          4.08\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"work_in_progress\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 235.44061859356583,\n        \"min\": 282.0,\n        \"max\": 1601.0,\n
```

```

{"num_unique_values": 24,\n
 "samples": [\n
  1436.0\n
 ],\n
 "semantic_type": "\"",\n
 "description": "\""\n
 }\n
 },\n
 {\n
 "column": "over_time",\n
 "properties": {\n
 "dtype": "number",\n
 "std": 2729,\n
 "min": 240,\n
 "max": 7080,\n
 "num_unique_values": 18,\n
 "samples": [\n
  6840\n
 ],\n
 "semantic_type": "\"",\n
 "description": "\""\n
 }\n
 },\n
 {\n
 "column": "incentive",\n
 "properties": {\n
 "dtype": "number",\n
 "std": 41,\n
 "min": 0,\n
 "max": 113,\n
 "num_unique_values": 12,\n
 "samples": [\n
  56\n
 ],\n
 "semantic_type": "\"",\n
 "description": "\""\n
 }\n
 },\n
 {\n
 "column": "idle_time",\n
 "properties": {\n
 "dtype": "number",\n
 "std": 0.0,\n
 "min": 0.0,\n
 "max": 0.0,\n
 "num_unique_values": 1,\n
 "samples": [\n
  0.0\n
 ],\n
 "semantic_type": "\"",\n
 "description": "\""\n
 }\n
 },\n
 {\n
 "column": "idle_men",\n
 "properties": {\n
 "dtype": "number",\n
 "std": 0,\n
 "min": 0,\n
 "max": 0,\n
 "num_unique_values": 1,\n
 "samples": [\n
  0\n
 ],\n
 "semantic_type": "\"",\n
 "description": "\""\n
 }\n
 },\n
 {\n
 "column": "no_of_style_change",\n
 "properties": {\n
 "dtype": "number",\n
 "std": 0,\n
 "min": 0,\n
 "max": 0,\n
 "num_unique_values": 1,\n
 "samples": [\n
  0\n
 ],\n
 "semantic_type": "\"",\n
 "description": "\""\n
 }\n
 },\n
 {\n
 "column": "no_of_workers",\n
 "properties": {\n
 "dtype": "number",\n
 "std": 23.20755007642304,\n
 "min": 2.0,\n
 "max": 59.0,\n
 "num_unique_values": 16,\n
 "samples": [\n
  57.0\n
 ],\n
 "semantic_type": "\"",\n
 "description": "\""\n
 }\n
 },\n
 {\n
 "column": "productivity_score",\n
 "properties": {\n
 "dtype": "number",\n
 "std": 18.385227641756785,\n
 "min": 28.698,\n
 "max": 100.046,\n
 "num_unique_values": 30,\n
 "samples": [\n
  60.071\n
 ],\n
 "semantic_type": "\"",\n
 "description": "\""\n
 }\n
 },\n
 {\n
 "column": "Date",\n
 "properties": {\n
 "dtype": "int32",\n
 "num_unique_values": 2,\n
 "samples": [\n
  31\n
 ],\n
 "semantic_type": "\"",\n
 "description": "\""\n
 }\n
 },\n
 {\n
 "column": "Month",\n
 "properties": {\n
 "dtype": "int32",\n
 "num_unique_values": 1,\n
 "samples": [\n
  1\n
 ],\n
 "semantic_type": "\"",\n
 "description": "\""\n
 }\n
 },\n
 {\n
 "column": "Year",\n
 "properties": {\n
 "dtype": "int32",\n
 "num_unique_values": 2,\n
 "samples": [\n
  2016\n
 ],\n
 "semantic_type": "\"",\n
 "description": "\""\n
 }\n
 }\n
 ]\n
 },\n
 "type": "dataframe"
}

```

```
# Drop Quarter5
df = df[df['quarter'] != 'Quarter5']
df['quarter'].unique()

array(['Quarter1', 'Quarter2', 'Quarter3', 'Quarter4'], dtype=object)
```

I decided to drop the Quarter5 data because, upon filtering, it showed that there are only 30 entries associated with Quarter5. This small number is not significant enough to justify creating a separate quarter, especially since it deviates from the standard four-quarter structure usually used in financial and performance analyses. Allocating these entries to Quarter1 could lead to misinterpretations of the data, as it may distort the understanding of seasonality and trends within the dataset.

```
# Check the unique entries
df['Team Code'].unique()

array([ 8,  1, 11, 12,  6,  7,  2,  3,  9, 10,  5,  4])

df['no_of_style_change'].unique()

array([0, 1, 2])

df['Month'].unique()

array([1, 2, 3], dtype=int32)

df['Year'].unique()

array([2015, 2016], dtype=int32)

cat_cols = ['quarter', 'day', 'Month', 'Year', 'Team Code',
            'no_of_style_change']
num_cols = ['std_minute_value', 'over_time', 'incentive', 'idle_time',
            'idle_men', 'no_of_workers', 'productivity_score', 'Date']
```

Change the "month," "year," "team code," and "no_of_style_change" columns into categorical types because the "month" column contains only values of 1, 2, or 3, representing three months; the "year" column consists only of the years 2015 and 2016; the "team code" ranges from 0 to 12; and the "no_of_style_change" column includes values of 0, 1, and 2.

```
from scipy.stats import skew, kurtosis

for col in num_cols:
    col_skew = skew(df[col].dropna())          # dropna to avoid issues
with NaNs
    col_kurt = kurtosis(df[col].dropna())      # by default, Fisher's
definition (normal ==> kurtosis 0)
    print(f"{col} - Skewness: {col_skew:.2f}, Kurtosis:
{col_kurt:.2f}")
kurtosis(df[col].dropna(), fisher=False)
```

```

std_minute_value - Skewness: 0.39, Kurtosis: -0.85
over_time - Skewness: 0.68, Kurtosis: 0.42
incentive - Skewness: 15.53, Kurtosis: 288.07
idle_time - Skewness: 20.14, Kurtosis: 424.41
idle_men - Skewness: 9.65, Kurtosis: 98.58
no_of_workers - Skewness: -0.12, Kurtosis: -1.78
productivity_score - Skewness: -0.81, Kurtosis: 0.36
Date - Skewness: 0.31, Kurtosis: -1.13

np.float64(1.8664251572797201)

```

Then, I also checked the skewness value for the numerical features to assess the distribution of the data. This helps determine whether the data is normally distributed or if it exhibits significant skewness, which could impact the choice of statistical methods and models used in the analysis. A skewness value of 20.14 (idle_time column) indicates a highly right-skewed distribution. This means that the majority of the data points are concentrated on the left side of the distribution. In practical terms, this high level of skewness can affect statistical analyses and modeling, making it necessary to consider data transformations or other approaches to better handle the distribution and ensure accurate insights.

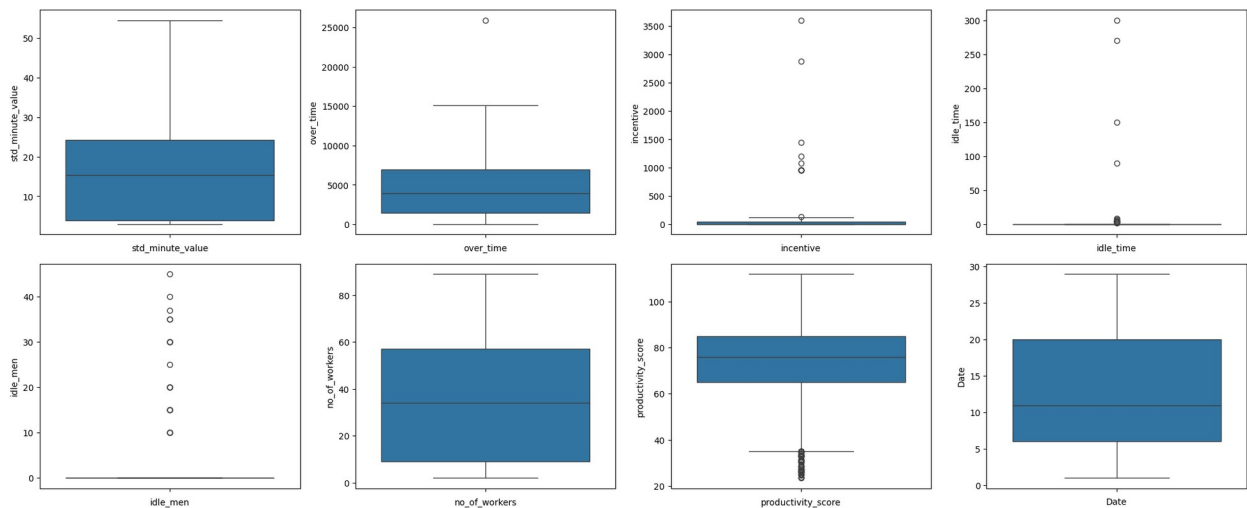
```

# Boxplot
fig = plt.figure(figsize=(20, 20))

for i, col in enumerate(num_cols, 1):
    plt.subplot(5, 4, i)
    sns.boxplot(df[col])
    plt.xlabel(col)

plt.tight_layout()
plt.show()

```



I decided to keep the outliers because removing them would result in a loss of valuable data. Since outliers can sometimes represent significant events or rare occurrences in the dataset,

preserving them allows for a more comprehensive analysis and helps ensure that important insights are not overlooked.

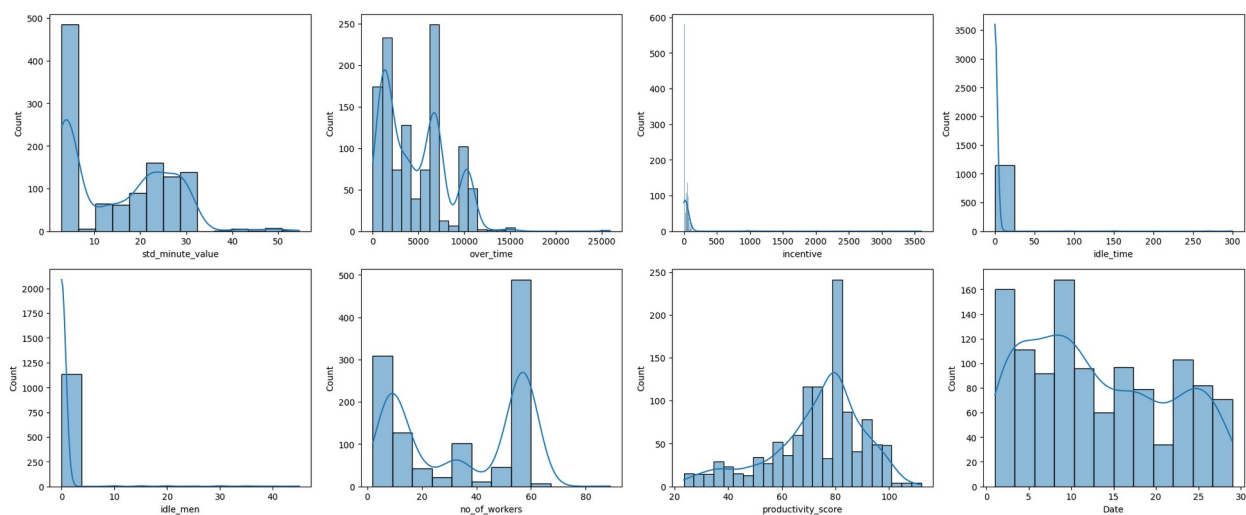
```
# Distribution of numerical features - Histogram
```

```
fig = plt.figure(figsize=(20, 20))
```

```
for i, col in enumerate(num_cols, 1):
    plt.subplot(5, 4, i)
    sns.histplot(df[col], kde=True)
    plt.xlabel(col)
```

```
plt.tight_layout()
```

```
plt.show()
```



The distributions of the numerical features indicate that the Standard Minute Value and Over Time features are left-skewed. In contrast, the Incentive, Idle Time, and Idle Men features are heavily peaked at zero. I chose not to handle these columns to avoid losing data. The Number of Workers displays a bimodal pattern, reflecting variability in team sizes, while the Productivity Score peaks around 80, suggesting that many teams perform well. Lastly, the Date distribution shows fluctuations, indicating varying levels of activity throughout the month.

```
# Barplot
```

```
# Create subplots
```

```
fig, axes = plt.subplots(nrows=3, ncols=2, figsize=(17, 8))
```

```
# Flatten the axes array for easier indexing
```

```
for i, ax in enumerate(axes.ravel()):
    if i < len(cat_cols): # Ensure we don't run into an error if
        cat_cols has fewer than 6 items
        avg_prod = df.groupby(cat_cols[i])
        ['productivity_score'].mean()
        avg_prod.plot.bar(ax=ax, color=sns.color_palette('Set2',
            n_colors=len(avg_prod)))
```

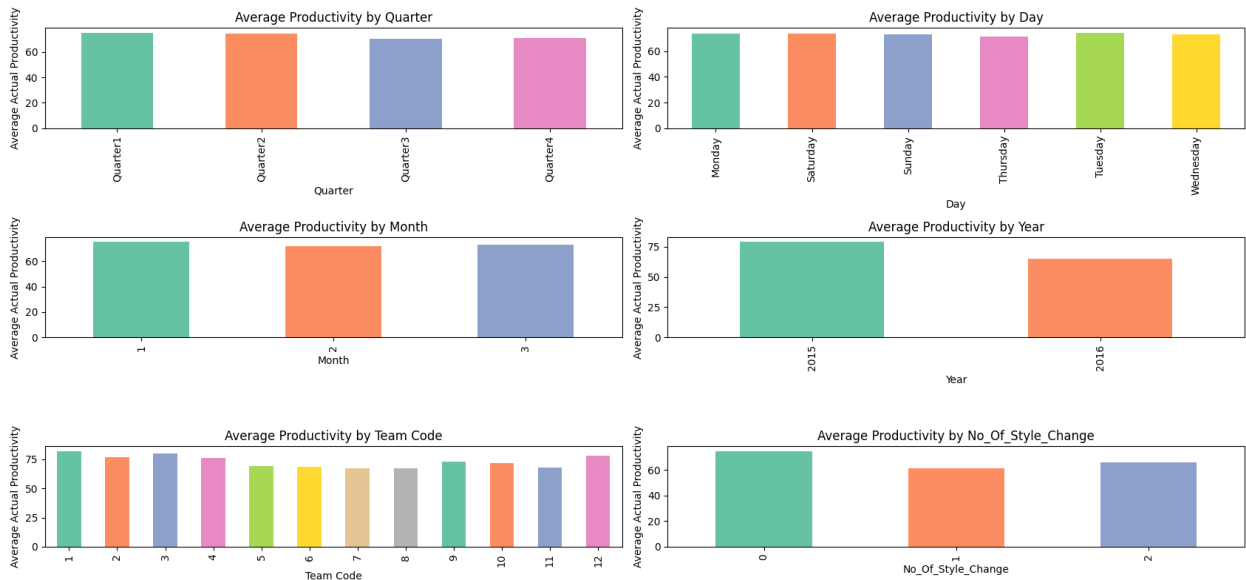


```

ax.set_title(f"Average Productivity by {cat_cols[i].title()}")
ax.set_xlabel(cat_cols[i].title())
ax.set_ylabel("Average Actual Productivity")
else:
    ax.axis('off') # Hide any unused subplots

# Adjust layout
fig.tight_layout()
plt.show()

```



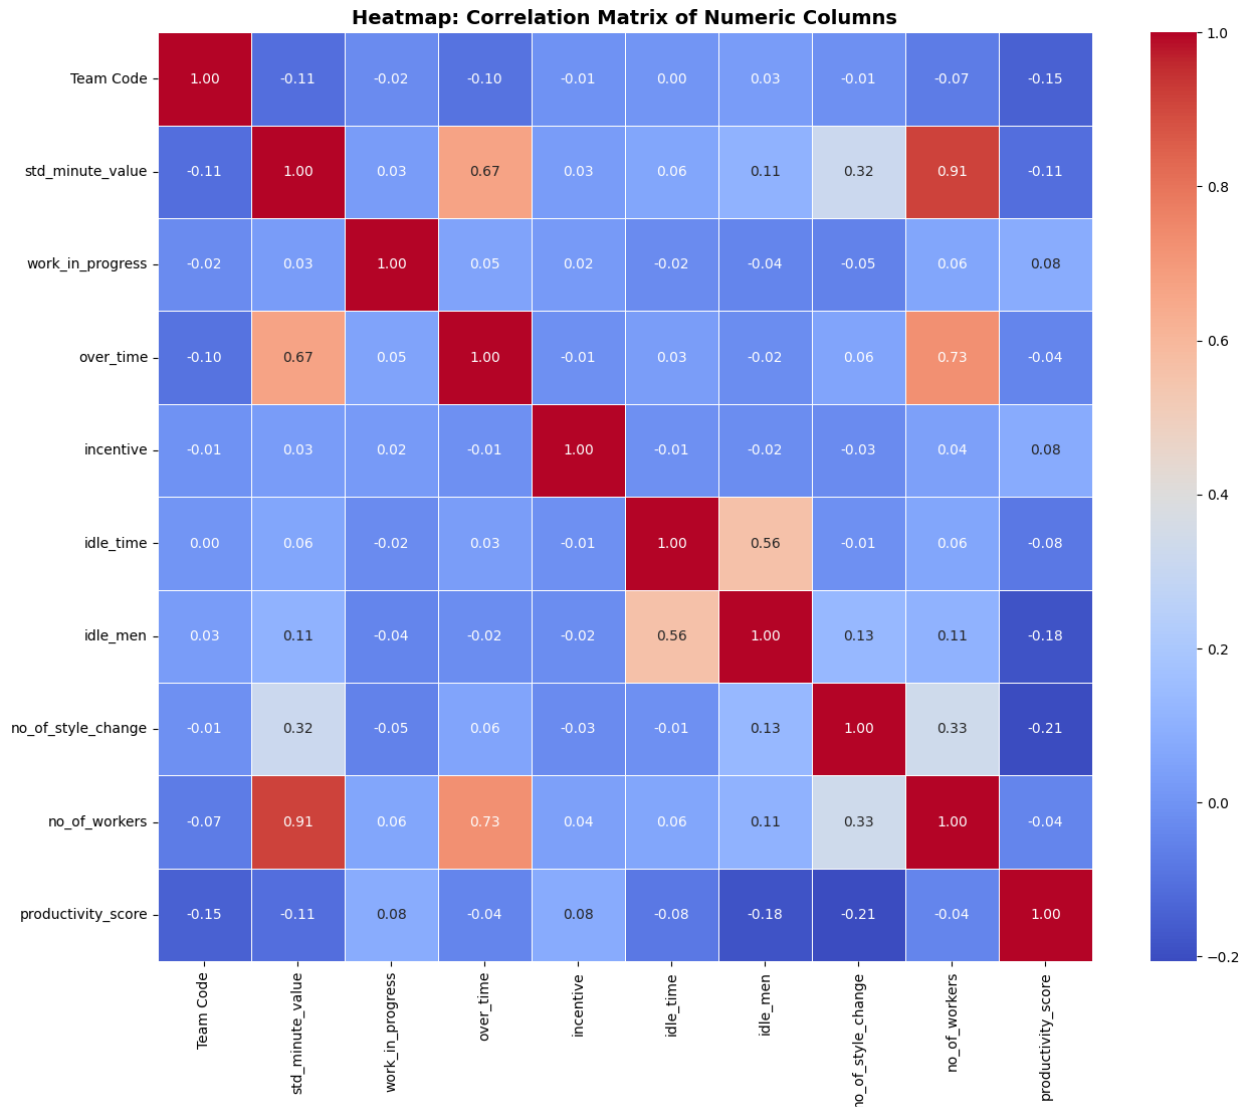
The count plots reveal important insights across various categorical features. In the **quarter** distribution, Quarter 1 has the highest entries, followed by Quarter 2, indicating a concentration of activities in the first half of the year. The **day** plot shows that Sunday and Thursday are the busiest, while Saturday and Monday have lower counts. For **months**, January and February are notably more active than March. In terms of **year**, 2015 has significantly more entries than 2016, suggesting a higher volume of data collection in that year. The **team code** counts are balanced, with Team 8 slightly leading, and the **number of style changes** overwhelmingly reflects no changes, indicating infrequent style alterations. Overall, these patterns offer valuable insights for operational and productivity strategies in the garment industry.

```

# Heatmap
numeric_columns = df.select_dtypes(include=['int64',
'float64']).columns
correlation_matrix = df[numeric_columns].corr()

plt.figure(figsize=(15, 12))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt='.2f', linewidths=0.5)
plt.title('Heatmap: Correlation Matrix of Numeric Columns',
fontsize=14, fontweight='bold')
plt.show()

```



From the heatmap, it shows that there are significant relationship among the features in garment dataset. Strong correlation between no_of_workers and std_minute_value (0.91) indicate that task require more time tend to involve larger teams. Negative Correlation such as Team Code and productivity score (-0.15) implies higher team assignment may relate to lower productivity levels

```
# delete date month year
df = df.drop(columns=['Date', 'Month', 'Year'])
```

Encoding

```
quarter_map = {'Quarter1': 1, 'Quarter2': 2, 'Quarter3': 3,
               'Quarter4': 4}
df['quarter'] = df['quarter'].map(quarter_map)
```

```

day_map = {
    'Monday': 0,
    'Tuesday': 1,
    'Wednesday': 2,
    'Thursday': 3,
    'Friday': 4,
    'Saturday': 5,
    'Sunday': 6,
}
df['day'] = df['day'].map(day_map)

```

Map the unique values in quarter and day columns to their corresponding numerical representations, this mapping method assign numerical values based on the unique categories present in each column.

```

encoder = OneHotEncoder(sparse_output=False, drop=None)
encoded_array = encoder.fit_transform(df[['Team Code']])

# Convert to integer type (if needed)
encoded_array = encoded_array.astype(int)

# Get column names
encoded_cols = encoder.get_feature_names_out(['Team Code'])

# Make it a DataFrame
df_encoded = pd.DataFrame(encoded_array, columns=encoded_cols,
index=df.index)

# Combine with original DataFrame
df_final = pd.concat([df.drop(columns=['Team Code']), df_encoded],
axis=1)

```

Then, I performed one-hot encoding for the Team Code because it contains discrete categorical values ranging from 1 to 12. One hot encoding effectively converts these categories into binary columns, which help to ensure the team code are treated as distinct categories, improving model's ability to learn patterns associated with each team.

Target Features

The goal is predicting productivity score each team in garment company.

Target Feature : "productivity_score"

```

# Split Target Features
x = df_final.drop(['productivity_score'], axis=1) # Drop the target
column from X
y = df_final['productivity_score'] # Target variable y

```

```
x.head()  
{ "type": "dataframe", "variable_name": "x" }
```

Since "productivity_score" is the target variable, it is removed from x and assigned to y for model training.

Split Data

Step 1 (test_size=0.3):

70% of the data is used for the train set.

30% of the data is used for the temporary set (x_temp, y_temp).

Step 2 (test_size=1/3):

From the 30% temporary data, 1/3 becomes the test set (which is 10% of the total data).

2/3 becomes the validation set (which is 20% of the total data).

So, with these two splits:

70% of the data is used for training.

20% of the data is used for validation.

10% of the data is used for testing.

```
# Split into train, validation, and test sets  
x_train, x_temp, y_train, y_temp = train_test_split(x, y,  
test_size=0.3, random_state=42)  
x_val, x_test, y_val, y_test = train_test_split(x_temp, y_temp,  
test_size=1/3, random_state=42)
```

Next, split the data into training, validation, and testing sets to ensure proper model evaluation. First, 80% of the data is allocated to the training set (x_train, y_train), while the remaining 20% is used as the test set (x_test, y_test). Then, 25% of the training set is further split into a validation set (x_val, y_val), ensuring a well-balanced dataset for effective model training, hyperparameter tuning, and final evaluation.

```
numeric_cols = x_train.select_dtypes(include=['int64',  
'float64']).columns  
categorical_cols = x_train.select_dtypes(include=['object']).columns
```

Scaling

```
scaler = StandardScaler()
x_train_scaled_num = scaler.fit_transform(x_train[numeric_cols])
x_val_scaled_num = scaler.transform(x_val[numeric_cols])
x_test_scaled_num = scaler.transform(x_test[numeric_cols])

x_train_scaled = pd.DataFrame(x_train_scaled_num,
                              columns=numeric_cols, index=x_train.index)
x_val_scaled = pd.DataFrame(x_val_scaled_num, columns=numeric_cols,
                             index=x_val.index)
x_test_scaled = pd.DataFrame(x_test_scaled_num, columns=numeric_cols,
                              index=x_test.index)

x_train_scaled = pd.concat([x_train_scaled,
                             x_train[categorical_cols]], axis=1)
x_val_scaled = pd.concat([x_val_scaled, x_val[categorical_cols]],
                           axis=1)
x_test_scaled = pd.concat([x_test_scaled, x_test[categorical_cols]],
                            axis=1)
```

In this project, I employed StandardScaler to preprocess the input features for the Artificial Neural Network (ANN) aimed at predicting productivity. The decision was based on several key factors:

1. **Handling Outliers:** The dataset contains various outliers, and using MinMaxScaler could compress the scale of non-outlier data, negatively impacting the model. StandardScaler, which normalizes features based on their mean and standard deviation, is less sensitive to extreme values.
2. **Addressing Skewed Distributions:** Many features are not uniformly distributed and exhibit skewness. MinMaxScaler assumes a uniform spread, which isn't suitable here, while StandardScaler effectively centers the data around zero and scales it according to its spread.
3. **Enhancing Neural Network Stability and Performance:** Standardization leads to faster and more reliable convergence for neural networks. By ensuring inputs have a mean of 0 and a standard deviation of 1, each neuron receives data on a consistent scale, resulting in more stable gradient updates during training.

With these considerations, StandardScaler is the optimal choice for scaling in this scenario, promoting robust model performance and efficient training.

Building Tensor Dataset

```
train_ds =
tf.data.Dataset.from_tensor_slices((x_train_scaled,y_train)).batch(32)
.shuffle(10)
```

```

test_ds =
tf.data.Dataset.from_tensor_slices((x_test_scaled,y_test)).batch(32)
val_ds =
tf.data.Dataset.from_tensor_slices((x_val_scaled,y_val)).batch(32)

# Print x_train shape
print(x_train_scaled.shape)

(807, 22)

train_ds

<_ShuffleDataset element_spec=(TensorSpec(shape=(None, 22),
dtype=tf.float64, name=None), TensorSpec(shape=(None, ),
dtype=tf.float64, name=None))>

```

Baseline Model

A baseline model is a starting point in a project. It is a simple model that helps us to know how well basic methods perform.

Sequential

```

model = tf.keras.Sequential([
    Dense(128, activation='relu', input_shape=(22,)), # Neuron minimal
2X from input
    Dense(64, activation='relu'),
    Dense(1, activation='linear')
])

```

The first model is a Sequential model with three hidden layers. The first layer has 128 neurons, followed by a second layer with 64 neurons, both using ReLU activation to capture complex patterns in the data.

```

model.compile(
    optimizer=tf.keras.optimizers.Adam(),
    loss='mean_squared_error',
    metrics=['mean_absolute_error', 'mean_squared_error']
)

```

```

# Model Summary
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape
--------------	--------------

Param #		
dense (Dense)	(None, 128)	
2,944		
dense_1 (Dense)	(None, 64)	
8,256		
dense_2 (Dense)	(None, 1)	
65		

Total params: 11,265 (44.00 KB)

Trainable params: 11,265 (44.00 KB)

Non-trainable params: 0 (0.00 B)

The minimum required number of epochs is 10, but in this case, I used 50 epochs to give the model more time to learn and potentially achieve better performance.

Train the model

```
history = model.fit(x_train_scaled, y_train,
                    validation_data=(x_val_scaled, y_val),
                    epochs=50,
                    batch_size=32,
                    verbose=1)
```

Epoch 1/50

```
26/26 ————— 5s 64ms/step - loss: 5358.6362 -
mean_absolute_error: 71.0948 - mean_squared_error: 5358.6362 -
val_loss: 4651.6113 - val_mean_absolute_error: 66.2169 -
val_mean_squared_error: 4651.6113
```

Epoch 2/50

```
26/26 ————— 1s 26ms/step - loss: 4678.4561 -
mean_absolute_error: 65.9596 - mean_squared_error: 4678.4561 -
val_loss: 3248.8472 - val_mean_absolute_error: 54.5327 -
val_mean_squared_error: 3248.8472
```

Epoch 3/50

```
26/26 ————— 0s 15ms/step - loss: 2986.9148 -
mean_absolute_error: 51.5474 - mean_squared_error: 2986.9148 -
val_loss: 1464.6555 - val_mean_absolute_error: 35.1155 -
val_mean_squared_error: 1464.6555
```

Epoch 4/50

```
26/26 ————— 1s 18ms/step - loss: 1168.5156 -
mean_absolute_error: 30.5262 - mean_squared_error: 1168.5156 -
```

```
val_loss: 415.9604 - val_mean_absolute_error: 17.6726 -  
val_mean_squared_error: 415.9604  
Epoch 5/50  
26/26 ————— 0s 14ms/step - loss: 357.2380 -  
mean_absolute_error: 15.6365 - mean_squared_error: 357.2380 -  
val_loss: 297.3177 - val_mean_absolute_error: 13.2375 -  
val_mean_squared_error: 297.3177  
Epoch 6/50  
26/26 ————— 1s 14ms/step - loss: 293.4005 -  
mean_absolute_error: 13.2418 - mean_squared_error: 293.4005 -  
val_loss: 272.6161 - val_mean_absolute_error: 12.5761 -  
val_mean_squared_error: 272.6161  
Epoch 7/50  
26/26 ————— 1s 13ms/step - loss: 292.6281 -  
mean_absolute_error: 13.2627 - mean_squared_error: 292.6281 -  
val_loss: 262.5137 - val_mean_absolute_error: 12.4080 -  
val_mean_squared_error: 262.5137  
Epoch 8/50  
26/26 ————— 1s 11ms/step - loss: 282.7032 -  
mean_absolute_error: 12.9848 - mean_squared_error: 282.7032 -  
val_loss: 258.5030 - val_mean_absolute_error: 12.1534 -  
val_mean_squared_error: 258.5030  
Epoch 9/50  
26/26 ————— 1s 20ms/step - loss: 247.2510 -  
mean_absolute_error: 12.0886 - mean_squared_error: 247.2510 -  
val_loss: 252.7126 - val_mean_absolute_error: 11.9085 -  
val_mean_squared_error: 252.7126  
Epoch 10/50  
26/26 ————— 0s 11ms/step - loss: 280.4243 -  
mean_absolute_error: 12.7915 - mean_squared_error: 280.4243 -  
val_loss: 251.5667 - val_mean_absolute_error: 11.8845 -  
val_mean_squared_error: 251.5667  
Epoch 11/50  
26/26 ————— 1s 13ms/step - loss: 256.1978 -  
mean_absolute_error: 12.3029 - mean_squared_error: 256.1978 -  
val_loss: 249.9779 - val_mean_absolute_error: 11.7245 -  
val_mean_squared_error: 249.9779  
Epoch 12/50  
26/26 ————— 1s 22ms/step - loss: 223.1900 -  
mean_absolute_error: 11.3838 - mean_squared_error: 223.1900 -  
val_loss: 246.1589 - val_mean_absolute_error: 11.6395 -  
val_mean_squared_error: 246.1589  
Epoch 13/50  
26/26 ————— 1s 20ms/step - loss: 233.5306 -  
mean_absolute_error: 11.5611 - mean_squared_error: 233.5306 -  
val_loss: 243.6043 - val_mean_absolute_error: 11.6043 -  
val_mean_squared_error: 243.6043  
Epoch 14/50  
26/26 ————— 1s 15ms/step - loss: 239.5105 -
```



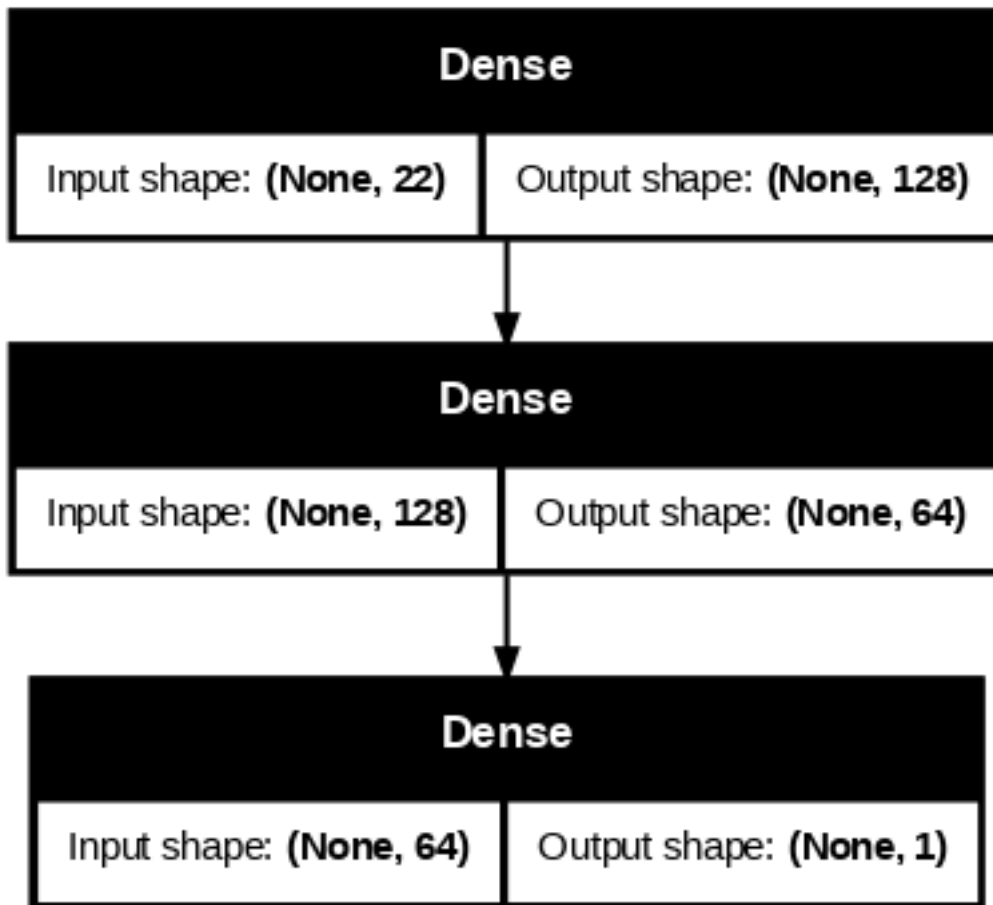
```
mean_absolute_error: 11.9045 - mean_squared_error: 239.5105 -  
val_loss: 241.5325 - val_mean_absolute_error: 11.5142 -  
val_mean_squared_error: 241.5325  
Epoch 15/50  
26/26 ————— 1s 11ms/step - loss: 225.6774 -  
mean_absolute_error: 11.5201 - mean_squared_error: 225.6774 -  
val_loss: 241.7192 - val_mean_absolute_error: 11.4719 -  
val_mean_squared_error: 241.7192  
Epoch 16/50  
26/26 ————— 1s 25ms/step - loss: 207.6405 -  
mean_absolute_error: 11.0020 - mean_squared_error: 207.6405 -  
val_loss: 239.9426 - val_mean_absolute_error: 11.4519 -  
val_mean_squared_error: 239.9426  
Epoch 17/50  
26/26 ————— 1s 31ms/step - loss: 231.7247 -  
mean_absolute_error: 11.7652 - mean_squared_error: 231.7247 -  
val_loss: 238.0210 - val_mean_absolute_error: 11.2842 -  
val_mean_squared_error: 238.0210  
Epoch 18/50  
26/26 ————— 1s 33ms/step - loss: 224.5269 -  
mean_absolute_error: 11.2000 - mean_squared_error: 224.5269 -  
val_loss: 236.8561 - val_mean_absolute_error: 11.3659 -  
val_mean_squared_error: 236.8561  
Epoch 19/50  
26/26 ————— 1s 19ms/step - loss: 228.3874 -  
mean_absolute_error: 11.6245 - mean_squared_error: 228.3874 -  
val_loss: 238.9326 - val_mean_absolute_error: 11.2855 -  
val_mean_squared_error: 238.9326  
Epoch 20/50  
26/26 ————— 1s 17ms/step - loss: 214.1136 -  
mean_absolute_error: 10.9041 - mean_squared_error: 214.1136 -  
val_loss: 236.8804 - val_mean_absolute_error: 11.2621 -  
val_mean_squared_error: 236.8804  
Epoch 21/50  
26/26 ————— 1s 16ms/step - loss: 212.7799 -  
mean_absolute_error: 10.8769 - mean_squared_error: 212.7799 -  
val_loss: 231.9994 - val_mean_absolute_error: 11.1343 -  
val_mean_squared_error: 231.9994  
Epoch 22/50  
26/26 ————— 0s 11ms/step - loss: 222.4274 -  
mean_absolute_error: 11.2558 - mean_squared_error: 222.4274 -  
val_loss: 233.1278 - val_mean_absolute_error: 11.1454 -  
val_mean_squared_error: 233.1278  
Epoch 23/50  
26/26 ————— 0s 14ms/step - loss: 214.4915 -  
mean_absolute_error: 11.0484 - mean_squared_error: 214.4915 -  
val_loss: 234.3152 - val_mean_absolute_error: 11.1391 -  
val_mean_squared_error: 234.3152  
Epoch 24/50
```

26/26 ————— 1s 11ms/step - loss: 212.6380 -
mean_absolute_error: 11.2341 - mean_squared_error: 212.6380 -
val_loss: 232.0361 - val_mean_absolute_error: 11.1735 -
val_mean_squared_error: 232.0361
Epoch 25/50
26/26 ————— 1s 23ms/step - loss: 174.7941 -
mean_absolute_error: 9.9115 - mean_squared_error: 174.7941 - val_loss:
230.3878 - val_mean_absolute_error: 11.1396 - val_mean_squared_error:
230.3878
Epoch 26/50
26/26 ————— 1s 6ms/step - loss: 226.4064 -
mean_absolute_error: 11.5877 - mean_squared_error: 226.4064 -
val_loss: 234.8536 - val_mean_absolute_error: 11.1456 -
val_mean_squared_error: 234.8536
Epoch 27/50
26/26 ————— 0s 6ms/step - loss: 214.7414 -
mean_absolute_error: 10.8571 - mean_squared_error: 214.7414 -
val_loss: 235.1540 - val_mean_absolute_error: 11.1974 -
val_mean_squared_error: 235.1540
Epoch 28/50
26/26 ————— 0s 5ms/step - loss: 205.0749 -
mean_absolute_error: 10.6132 - mean_squared_error: 205.0749 -
val_loss: 232.0501 - val_mean_absolute_error: 11.1346 -
val_mean_squared_error: 232.0501
Epoch 29/50
26/26 ————— 0s 6ms/step - loss: 202.8737 -
mean_absolute_error: 10.6755 - mean_squared_error: 202.8737 -
val_loss: 232.5561 - val_mean_absolute_error: 11.1213 -
val_mean_squared_error: 232.5561
Epoch 30/50
26/26 ————— 0s 5ms/step - loss: 194.5984 -
mean_absolute_error: 10.4603 - mean_squared_error: 194.5984 -
val_loss: 233.1612 - val_mean_absolute_error: 11.2009 -
val_mean_squared_error: 233.1612
Epoch 31/50
26/26 ————— 0s 5ms/step - loss: 206.9041 -
mean_absolute_error: 10.8685 - mean_squared_error: 206.9041 -
val_loss: 232.9809 - val_mean_absolute_error: 11.1372 -
val_mean_squared_error: 232.9809
Epoch 32/50
26/26 ————— 0s 5ms/step - loss: 194.3673 -
mean_absolute_error: 10.6117 - mean_squared_error: 194.3673 -
val_loss: 232.1484 - val_mean_absolute_error: 11.0373 -
val_mean_squared_error: 232.1484
Epoch 33/50
26/26 ————— 0s 6ms/step - loss: 199.3568 -
mean_absolute_error: 10.4863 - mean_squared_error: 199.3568 -
val_loss: 232.0366 - val_mean_absolute_error: 11.3342 -
val_mean_squared_error: 232.0366

Epoch 34/50
26/26 ————— 0s 6ms/step - loss: 209.0277 -
mean_absolute_error: 10.9494 - mean_squared_error: 209.0277 -
val_loss: 232.1294 - val_mean_absolute_error: 11.0627 -
val_mean_squared_error: 232.1294
Epoch 35/50
26/26 ————— 0s 5ms/step - loss: 171.7456 -
mean_absolute_error: 9.7488 - mean_squared_error: 171.7456 - val_loss:
231.6498 - val_mean_absolute_error: 11.1930 - val_mean_squared_error:
231.6498
Epoch 36/50
26/26 ————— 0s 6ms/step - loss: 183.6318 -
mean_absolute_error: 10.3989 - mean_squared_error: 183.6318 -
val_loss: 231.7758 - val_mean_absolute_error: 11.1517 -
val_mean_squared_error: 231.7758
Epoch 37/50
26/26 ————— 0s 7ms/step - loss: 184.4121 -
mean_absolute_error: 10.2515 - mean_squared_error: 184.4121 -
val_loss: 230.2932 - val_mean_absolute_error: 11.1040 -
val_mean_squared_error: 230.2932
Epoch 38/50
26/26 ————— 0s 6ms/step - loss: 209.8366 -
mean_absolute_error: 10.8548 - mean_squared_error: 209.8366 -
val_loss: 229.6699 - val_mean_absolute_error: 10.9702 -
val_mean_squared_error: 229.6699
Epoch 39/50
26/26 ————— 0s 7ms/step - loss: 193.0451 -
mean_absolute_error: 10.4089 - mean_squared_error: 193.0451 -
val_loss: 231.5444 - val_mean_absolute_error: 11.1093 -
val_mean_squared_error: 231.5444
Epoch 40/50
26/26 ————— 0s 6ms/step - loss: 207.1039 -
mean_absolute_error: 10.6506 - mean_squared_error: 207.1039 -
val_loss: 234.1166 - val_mean_absolute_error: 11.1125 -
val_mean_squared_error: 234.1166
Epoch 41/50
26/26 ————— 0s 6ms/step - loss: 176.0679 -
mean_absolute_error: 9.9634 - mean_squared_error: 176.0679 - val_loss:
232.9312 - val_mean_absolute_error: 11.0181 - val_mean_squared_error:
232.9312
Epoch 42/50
26/26 ————— 0s 6ms/step - loss: 182.3673 -
mean_absolute_error: 10.0537 - mean_squared_error: 182.3673 -
val_loss: 231.9330 - val_mean_absolute_error: 11.0332 -
val_mean_squared_error: 231.9330
Epoch 43/50
26/26 ————— 0s 6ms/step - loss: 173.7043 -
mean_absolute_error: 9.7998 - mean_squared_error: 173.7043 - val_loss:
227.8138 - val_mean_absolute_error: 11.0959 - val_mean_squared_error:

```
227.8138
Epoch 44/50
26/26 _____ 0s 5ms/step - loss: 178.6293 -
mean_absolute_error: 10.0558 - mean_squared_error: 178.6293 -
val_loss: 232.7407 - val_mean_absolute_error: 10.9020 -
val_mean_squared_error: 232.7407
Epoch 45/50
26/26 _____ 0s 6ms/step - loss: 199.9169 -
mean_absolute_error: 10.4984 - mean_squared_error: 199.9169 -
val_loss: 231.7280 - val_mean_absolute_error: 11.0270 -
val_mean_squared_error: 231.7280
Epoch 46/50
26/26 _____ 0s 6ms/step - loss: 168.5697 -
mean_absolute_error: 9.6296 - mean_squared_error: 168.5697 - val_loss:
232.0538 - val_mean_absolute_error: 11.2452 - val_mean_squared_error:
232.0538
Epoch 47/50
26/26 _____ 0s 7ms/step - loss: 199.2104 -
mean_absolute_error: 10.7417 - mean_squared_error: 199.2104 -
val_loss: 234.8258 - val_mean_absolute_error: 10.9460 -
val_mean_squared_error: 234.8258
Epoch 48/50
26/26 _____ 0s 8ms/step - loss: 175.0897 -
mean_absolute_error: 9.9881 - mean_squared_error: 175.0897 - val_loss:
230.7715 - val_mean_absolute_error: 11.1372 - val_mean_squared_error:
230.7715
Epoch 49/50
26/26 _____ 0s 10ms/step - loss: 185.4055 -
mean_absolute_error: 10.2561 - mean_squared_error: 185.4055 -
val_loss: 234.3742 - val_mean_absolute_error: 10.9478 -
val_mean_squared_error: 234.3742
Epoch 50/50
26/26 _____ 0s 9ms/step - loss: 174.3193 -
mean_absolute_error: 9.7612 - mean_squared_error: 174.3193 - val_loss:
232.0248 - val_mean_absolute_error: 10.9970 - val_mean_squared_error:
232.0248

plot_model(model, show_shapes=True, dpi=75)
```

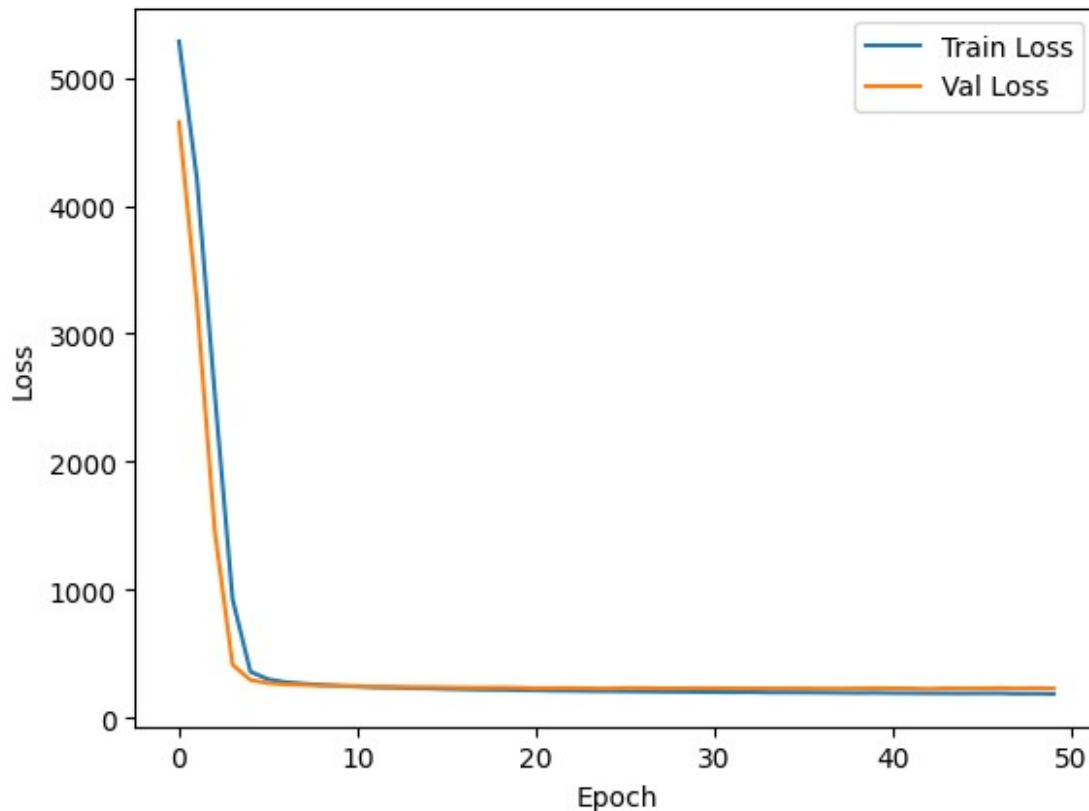


```
# Prediction
y_pred = model.predict(x_test_scaled)

4/4 ————— 0s 33ms/step
```

Evaluate

```
train_loss = history.history['loss']
val_loss = history.history['val_loss']
plt.plot(train_loss, label="Train Loss")
plt.plot(val_loss, label='Val Loss')
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



The loss curve shows that both training and validation loss dropped significantly in the first few epochs and then gradually stabilized, indicating the model learned quickly without overfitting. The small gap between train and validation loss suggests good generalization. However, since the loss decreased very rapidly, it might also indicate that the model is too simple for the dataset and may not be capturing more complex patterns.

3 Metric Analysis

```
R2 = r2_score(y_test, y_pred)
MAE = mean_absolute_error(y_test, y_pred)
MSE = mean_squared_error(y_test, y_pred)
```

```
print("R² Score=", R2)
print("Mean Absolute Error=", MAE)
print("Mean Squared Error=", MSE)
```

```
R² Score= 0.12921236289162408
Mean Absolute Error= 12.86850448937252
Mean Squared Error= 308.3501299756072
```

Plot Mean Absolute Error (MAE)

```
plt.figure(figsize=(12, 4))
```

```
plt.subplot(1, 2, 1)
```

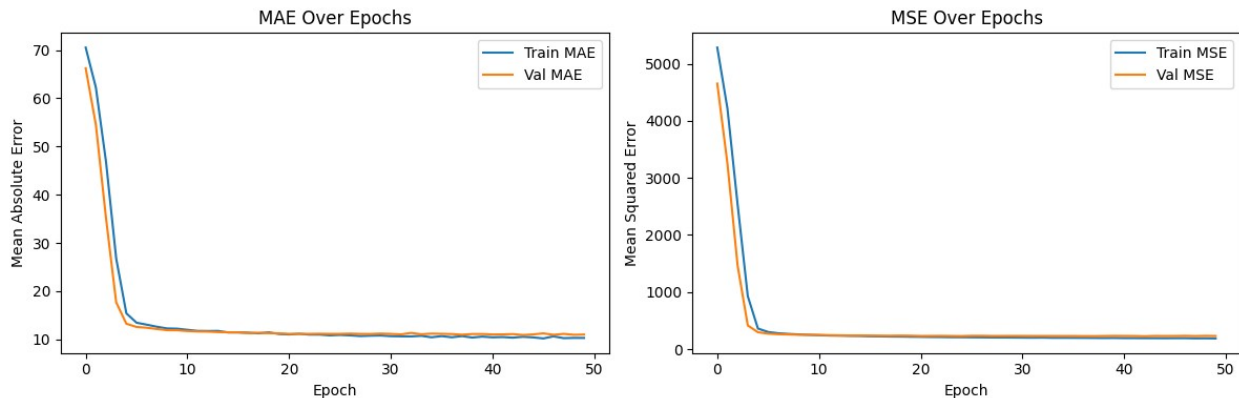
```
plt.plot(history.history['mean_absolute_error'], label='Train MAE')
```

```
plt.plot(history.history['val_mean_absolute_error'], label='Val MAE')
```

```
plt.xlabel('Epoch')
plt.ylabel('Mean Absolute Error')
plt.title('MAE Over Epochs')
plt.legend()

# Plot Mean Squared Error (MSE)
plt.subplot(1, 2, 2)
plt.plot(history.history['mean_squared_error'], label='Train MSE')
plt.plot(history.history['val_mean_squared_error'], label='Val MSE')
plt.xlabel('Epoch')
plt.ylabel('Mean Squared Error')
plt.title('MSE Over Epochs')
plt.legend()

plt.tight_layout()
plt.show()
```



The model learned quickly, as shown by the loss curves stabilizing without overfitting. However, the final metrics are

1. $R^2 = 0.129$,
2. MAE = 12.87, and
3. MSE = 308.35

indicate that the model's predictive power is low, explaining only 12.9% of the variance. This suggests the model is too simple for the dataset, and improvements are needed. High MAE and MSE value also indicate that the still need more modification to predict accurately.

Functional API

```
from tensorflow.keras.layers import Input

inputs = Input(shape=(x_train_scaled.shape[1],))

x = Dense(128, activation='relu')(inputs)
x = Dense(64, activation='relu')(x)
outputs = Dense(1, activation='linear')(x)
```

The second model is built using the Functional API with three layers. The first layer has 128 neurons with ReLU activation, followed by a second layer with 64 neurons, also using ReLU activation to capture more complex patterns.

```
model_1 = Model(inputs=inputs, outputs=outputs)

model_1.compile(
    optimizer=tf.keras.optimizers.Adam(),
    loss='mean_squared_error',
    metrics=['mean_absolute_error', 'mean_squared_error']
)
```

```
# Model Summary
```

```
model_1.summary()
```

```
Model: "functional_1"
```

Layer (type)	Output Shape	
Param #		
input_layer_1 (InputLayer)	(None, 22)	
0		
dense_3 (Dense)	(None, 128)	
2,944		
dense_4 (Dense)	(None, 64)	
8,256		
dense_5 (Dense)	(None, 1)	
65		

```
Total params: 11,265 (44.00 KB)
```

```
Trainable params: 11,265 (44.00 KB)
```

```
Non-trainable params: 0 (0.00 B)
```

The minimum required number of epochs in this model is also 10, but in this case, I also used 50 epochs to give the model more time to learn and potentially achieve better performance.

```
# Train the model
```

```
history = model_1.fit(x_train_scaled, y_train,
```



```
validation_data=(x_val_scaled, y_val),  
epochs=50,  
batch_size=32,  
verbose=1)
```

Epoch 1/50

```
26/26 _____ 2s 14ms/step - loss: 5449.6519 -  
mean_absolute_error: 71.6416 - mean_squared_error: 5449.6519 -  
val_loss: 4719.4854 - val_mean_absolute_error: 66.7400 -  
val_mean_squared_error: 4719.4854
```

Epoch 2/50

```
26/26 _____ 0s 7ms/step - loss: 4672.2324 -  
mean_absolute_error: 66.0836 - mean_squared_error: 4672.2324 -  
val_loss: 3360.1025 - val_mean_absolute_error: 55.5893 -  
val_mean_squared_error: 3360.1025
```

Epoch 3/50

```
26/26 _____ 0s 8ms/step - loss: 2949.0913 -  
mean_absolute_error: 51.2472 - mean_squared_error: 2949.0913 -  
val_loss: 1486.9247 - val_mean_absolute_error: 35.4299 -  
val_mean_squared_error: 1486.9247
```

Epoch 4/50

```
26/26 _____ 0s 7ms/step - loss: 1160.7073 -  
mean_absolute_error: 30.6357 - mean_squared_error: 1160.7073 -  
val_loss: 388.8639 - val_mean_absolute_error: 16.6818 -  
val_mean_squared_error: 388.8639
```

Epoch 5/50

```
26/26 _____ 0s 6ms/step - loss: 383.6802 -  
mean_absolute_error: 15.9054 - mean_squared_error: 383.6802 -  
val_loss: 288.8310 - val_mean_absolute_error: 12.9639 -  
val_mean_squared_error: 288.8310
```

Epoch 6/50

```
26/26 _____ 0s 7ms/step - loss: 281.7173 -  
mean_absolute_error: 12.8925 - mean_squared_error: 281.7173 -  
val_loss: 272.9529 - val_mean_absolute_error: 12.5817 -  
val_mean_squared_error: 272.9529
```

Epoch 7/50

```
26/26 _____ 0s 8ms/step - loss: 302.8765 -  
mean_absolute_error: 13.4368 - mean_squared_error: 302.8765 -  
val_loss: 261.6469 - val_mean_absolute_error: 12.4264 -  
val_mean_squared_error: 261.6469
```

Epoch 8/50

```
26/26 _____ 0s 6ms/step - loss: 278.5079 -  
mean_absolute_error: 13.2208 - mean_squared_error: 278.5079 -  
val_loss: 257.1906 - val_mean_absolute_error: 12.2425 -  
val_mean_squared_error: 257.1906
```

Epoch 9/50

```
26/26 _____ 0s 7ms/step - loss: 255.0793 -  
mean_absolute_error: 12.3952 - mean_squared_error: 255.0793 -  
val_loss: 252.3163 - val_mean_absolute_error: 11.9487 -  
val_mean_squared_error: 252.3163
```

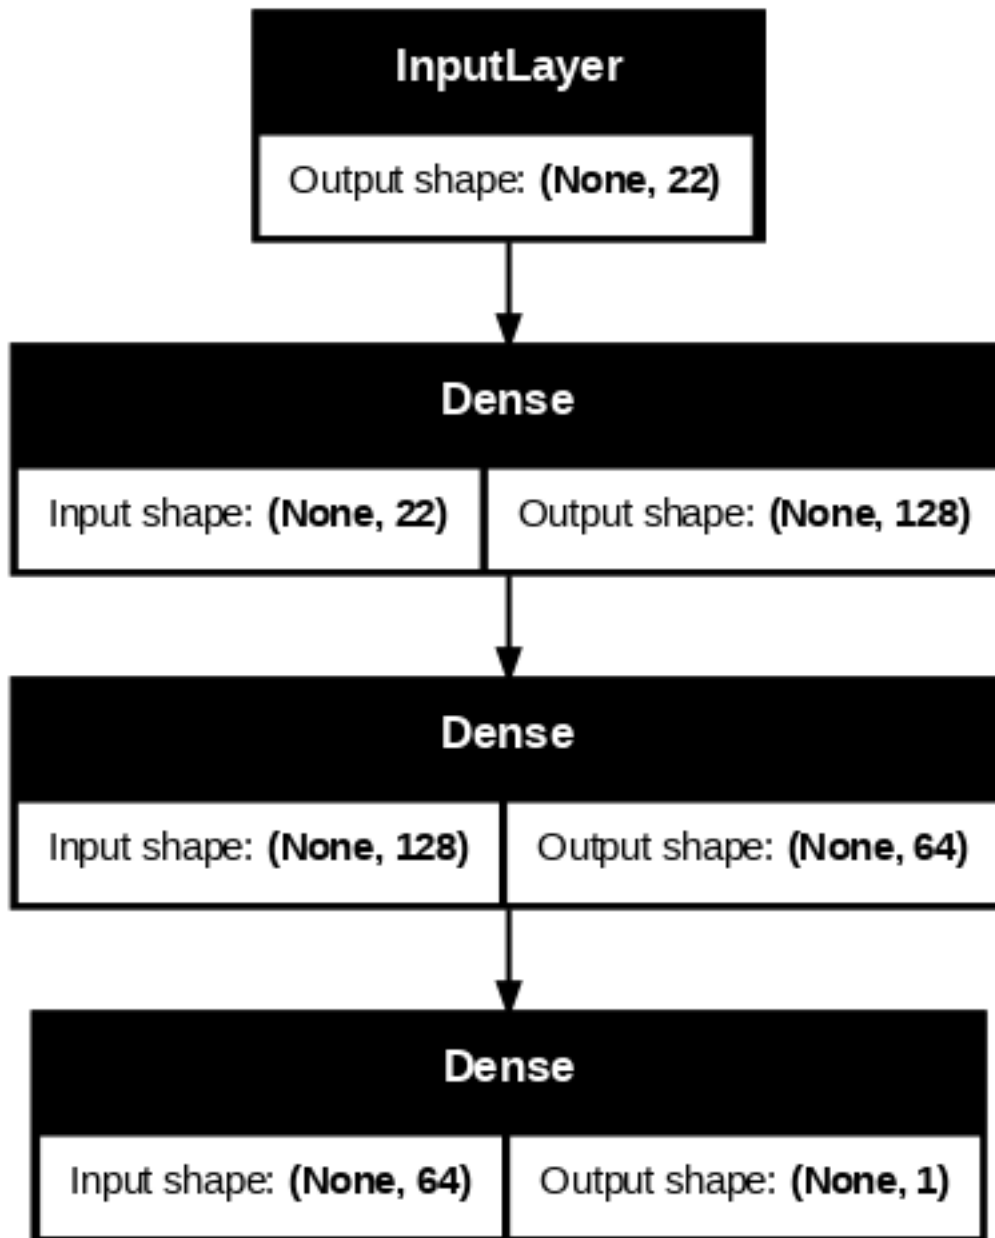
Epoch 10/50
26/26 ————— 0s 7ms/step - loss: 268.9190 -
mean_absolute_error: 12.8541 - mean_squared_error: 268.9190 -
val_loss: 251.0285 - val_mean_absolute_error: 11.8848 -
val_mean_squared_error: 251.0285
Epoch 11/50
26/26 ————— 0s 7ms/step - loss: 253.7633 -
mean_absolute_error: 12.1895 - mean_squared_error: 253.7633 -
val_loss: 249.3957 - val_mean_absolute_error: 11.7386 -
val_mean_squared_error: 249.3957
Epoch 12/50
26/26 ————— 0s 6ms/step - loss: 215.2738 -
mean_absolute_error: 11.3137 - mean_squared_error: 215.2738 -
val_loss: 244.9789 - val_mean_absolute_error: 11.5944 -
val_mean_squared_error: 244.9789
Epoch 13/50
26/26 ————— 0s 7ms/step - loss: 239.7575 -
mean_absolute_error: 11.9111 - mean_squared_error: 239.7575 -
val_loss: 244.9755 - val_mean_absolute_error: 11.5975 -
val_mean_squared_error: 244.9755
Epoch 14/50
26/26 ————— 0s 7ms/step - loss: 226.6929 -
mean_absolute_error: 11.3882 - mean_squared_error: 226.6929 -
val_loss: 244.3934 - val_mean_absolute_error: 11.4206 -
val_mean_squared_error: 244.3934
Epoch 15/50
26/26 ————— 0s 7ms/step - loss: 212.6314 -
mean_absolute_error: 11.0316 - mean_squared_error: 212.6314 -
val_loss: 241.6690 - val_mean_absolute_error: 11.3856 -
val_mean_squared_error: 241.6690
Epoch 16/50
26/26 ————— 0s 7ms/step - loss: 223.6355 -
mean_absolute_error: 11.5694 - mean_squared_error: 223.6355 -
val_loss: 238.4850 - val_mean_absolute_error: 11.2654 -
val_mean_squared_error: 238.4850
Epoch 17/50
26/26 ————— 0s 8ms/step - loss: 221.3133 -
mean_absolute_error: 11.2275 - mean_squared_error: 221.3133 -
val_loss: 237.2972 - val_mean_absolute_error: 11.2516 -
val_mean_squared_error: 237.2972
Epoch 18/50
26/26 ————— 0s 6ms/step - loss: 221.4541 -
mean_absolute_error: 11.3452 - mean_squared_error: 221.4541 -
val_loss: 238.7753 - val_mean_absolute_error: 11.3478 -
val_mean_squared_error: 238.7753
Epoch 19/50
26/26 ————— 0s 6ms/step - loss: 211.2988 -
mean_absolute_error: 11.1453 - mean_squared_error: 211.2988 -
val_loss: 235.6654 - val_mean_absolute_error: 11.2473 -

```
val_mean_squared_error: 235.6654
Epoch 20/50
26/26 _____ 0s 5ms/step - loss: 206.2696 -
mean_absolute_error: 10.9548 - mean_squared_error: 206.2696 -
val_loss: 236.5964 - val_mean_absolute_error: 11.0988 -
val_mean_squared_error: 236.5964
Epoch 21/50
26/26 _____ 0s 5ms/step - loss: 197.1338 -
mean_absolute_error: 10.8700 - mean_squared_error: 197.1338 -
val_loss: 236.8905 - val_mean_absolute_error: 11.2747 -
val_mean_squared_error: 236.8905
Epoch 22/50
26/26 _____ 0s 5ms/step - loss: 200.4566 -
mean_absolute_error: 10.9924 - mean_squared_error: 200.4566 -
val_loss: 237.7415 - val_mean_absolute_error: 11.2063 -
val_mean_squared_error: 237.7415
Epoch 23/50
26/26 _____ 0s 7ms/step - loss: 229.6906 -
mean_absolute_error: 11.6454 - mean_squared_error: 229.6906 -
val_loss: 238.9200 - val_mean_absolute_error: 11.2707 -
val_mean_squared_error: 238.9200
Epoch 24/50
26/26 _____ 0s 7ms/step - loss: 187.5894 -
mean_absolute_error: 10.3850 - mean_squared_error: 187.5894 -
val_loss: 236.9964 - val_mean_absolute_error: 11.3057 -
val_mean_squared_error: 236.9964
Epoch 25/50
26/26 _____ 0s 7ms/step - loss: 192.1990 -
mean_absolute_error: 10.5181 - mean_squared_error: 192.1990 -
val_loss: 235.8731 - val_mean_absolute_error: 11.2080 -
val_mean_squared_error: 235.8731
Epoch 26/50
26/26 _____ 0s 6ms/step - loss: 201.0883 -
mean_absolute_error: 10.8371 - mean_squared_error: 201.0883 -
val_loss: 235.9014 - val_mean_absolute_error: 11.1982 -
val_mean_squared_error: 235.9014
Epoch 27/50
26/26 _____ 0s 6ms/step - loss: 214.3575 -
mean_absolute_error: 10.9688 - mean_squared_error: 214.3575 -
val_loss: 234.0916 - val_mean_absolute_error: 11.0751 -
val_mean_squared_error: 234.0916
Epoch 28/50
26/26 _____ 0s 5ms/step - loss: 183.3612 -
mean_absolute_error: 10.2086 - mean_squared_error: 183.3612 -
val_loss: 235.3814 - val_mean_absolute_error: 11.1152 -
val_mean_squared_error: 235.3814
Epoch 29/50
26/26 _____ 0s 5ms/step - loss: 185.6768 -
mean_absolute_error: 10.2622 - mean_squared_error: 185.6768 -
```

```
val_loss: 235.7099 - val_mean_absolute_error: 11.1809 -  
val_mean_squared_error: 235.7099  
Epoch 30/50  
26/26 ━━━━━━━━━━━ 0s 6ms/step - loss: 198.3161 -  
mean_absolute_error: 10.6166 - mean_squared_error: 198.3161 -  
val_loss: 236.4979 - val_mean_absolute_error: 11.2212 -  
val_mean_squared_error: 236.4979  
Epoch 31/50  
26/26 ━━━━━━━━━━━ 0s 6ms/step - loss: 196.0229 -  
mean_absolute_error: 10.4616 - mean_squared_error: 196.0229 -  
val_loss: 235.2513 - val_mean_absolute_error: 11.1368 -  
val_mean_squared_error: 235.2513  
Epoch 32/50  
26/26 ━━━━━━━━━━━ 0s 11ms/step - loss: 189.6390 -  
mean_absolute_error: 10.4473 - mean_squared_error: 189.6390 -  
val_loss: 235.2974 - val_mean_absolute_error: 11.0936 -  
val_mean_squared_error: 235.2974  
Epoch 33/50  
26/26 ━━━━━━━━━━━ 1s 8ms/step - loss: 192.6635 -  
mean_absolute_error: 10.4926 - mean_squared_error: 192.6635 -  
val_loss: 235.0291 - val_mean_absolute_error: 11.1484 -  
val_mean_squared_error: 235.0291  
Epoch 34/50  
26/26 ━━━━━━━━━━━ 0s 9ms/step - loss: 191.3441 -  
mean_absolute_error: 10.5039 - mean_squared_error: 191.3441 -  
val_loss: 236.7380 - val_mean_absolute_error: 11.1253 -  
val_mean_squared_error: 236.7380  
Epoch 35/50  
26/26 ━━━━━━━━━━━ 0s 8ms/step - loss: 181.9235 -  
mean_absolute_error: 10.0873 - mean_squared_error: 181.9235 -  
val_loss: 236.2333 - val_mean_absolute_error: 11.1987 -  
val_mean_squared_error: 236.2333  
Epoch 36/50  
26/26 ━━━━━━━━━━━ 0s 9ms/step - loss: 205.1174 -  
mean_absolute_error: 11.0868 - mean_squared_error: 205.1174 -  
val_loss: 236.7294 - val_mean_absolute_error: 11.0927 -  
val_mean_squared_error: 236.7294  
Epoch 37/50  
26/26 ━━━━━━━━━━━ 0s 10ms/step - loss: 193.8205 -  
mean_absolute_error: 10.5269 - mean_squared_error: 193.8205 -  
val_loss: 236.9320 - val_mean_absolute_error: 11.1058 -  
val_mean_squared_error: 236.9320  
Epoch 38/50  
26/26 ━━━━━━━━━━━ 0s 7ms/step - loss: 187.4594 -  
mean_absolute_error: 10.1406 - mean_squared_error: 187.4594 -  
val_loss: 238.5517 - val_mean_absolute_error: 11.1805 -  
val_mean_squared_error: 238.5517  
Epoch 39/50  
26/26 ━━━━━━━━━━━ 0s 5ms/step - loss: 202.6889 -
```

```
mean_absolute_error: 10.6678 - mean_squared_error: 202.6889 -  
val_loss: 236.6911 - val_mean_absolute_error: 11.1205 -  
val_mean_squared_error: 236.6911  
Epoch 40/50  
26/26 ————— 0s 5ms/step - loss: 221.0382 -  
mean_absolute_error: 11.2026 - mean_squared_error: 221.0382 -  
val_loss: 239.5552 - val_mean_absolute_error: 11.1308 -  
val_mean_squared_error: 239.5552  
Epoch 41/50  
26/26 ————— 0s 5ms/step - loss: 193.4230 -  
mean_absolute_error: 10.4609 - mean_squared_error: 193.4230 -  
val_loss: 237.8185 - val_mean_absolute_error: 11.1688 -  
val_mean_squared_error: 237.8185  
Epoch 42/50  
26/26 ————— 0s 5ms/step - loss: 186.9290 -  
mean_absolute_error: 10.1775 - mean_squared_error: 186.9290 -  
val_loss: 236.7873 - val_mean_absolute_error: 11.2795 -  
val_mean_squared_error: 236.7873  
Epoch 43/50  
26/26 ————— 0s 5ms/step - loss: 177.7643 -  
mean_absolute_error: 10.2073 - mean_squared_error: 177.7643 -  
val_loss: 237.3213 - val_mean_absolute_error: 11.0702 -  
val_mean_squared_error: 237.3213  
Epoch 44/50  
26/26 ————— 0s 8ms/step - loss: 184.3974 -  
mean_absolute_error: 10.2182 - mean_squared_error: 184.3974 -  
val_loss: 237.2388 - val_mean_absolute_error: 11.0781 -  
val_mean_squared_error: 237.2388  
Epoch 45/50  
26/26 ————— 0s 7ms/step - loss: 177.1453 -  
mean_absolute_error: 10.0129 - mean_squared_error: 177.1453 -  
val_loss: 240.5054 - val_mean_absolute_error: 11.2322 -  
val_mean_squared_error: 240.5054  
Epoch 46/50  
26/26 ————— 0s 6ms/step - loss: 183.9633 -  
mean_absolute_error: 10.0597 - mean_squared_error: 183.9633 -  
val_loss: 239.3889 - val_mean_absolute_error: 11.1797 -  
val_mean_squared_error: 239.3889  
Epoch 47/50  
26/26 ————— 0s 7ms/step - loss: 185.9175 -  
mean_absolute_error: 10.1802 - mean_squared_error: 185.9175 -  
val_loss: 240.2192 - val_mean_absolute_error: 11.0950 -  
val_mean_squared_error: 240.2192  
Epoch 48/50  
26/26 ————— 0s 5ms/step - loss: 186.1756 -  
mean_absolute_error: 10.1265 - mean_squared_error: 186.1756 -  
val_loss: 237.3077 - val_mean_absolute_error: 11.1763 -  
val_mean_squared_error: 237.3077  
Epoch 49/50
```

```
26/26 _____ 0s 7ms/step - loss: 196.4990 -  
mean_absolute_error: 10.6038 - mean_squared_error: 196.4990 -  
val_loss: 239.4495 - val_mean_absolute_error: 11.1059 -  
val_mean_squared_error: 239.4495  
Epoch 50/50  
26/26 _____ 0s 5ms/step - loss: 167.2845 -  
mean_absolute_error: 9.6527 - mean_squared_error: 167.2845 - val_loss:  
240.3244 - val_mean_absolute_error: 11.1592 - val_mean_squared_error:  
240.3244  
  
plot_model(model_1, show_shapes=True, dpi=75)
```

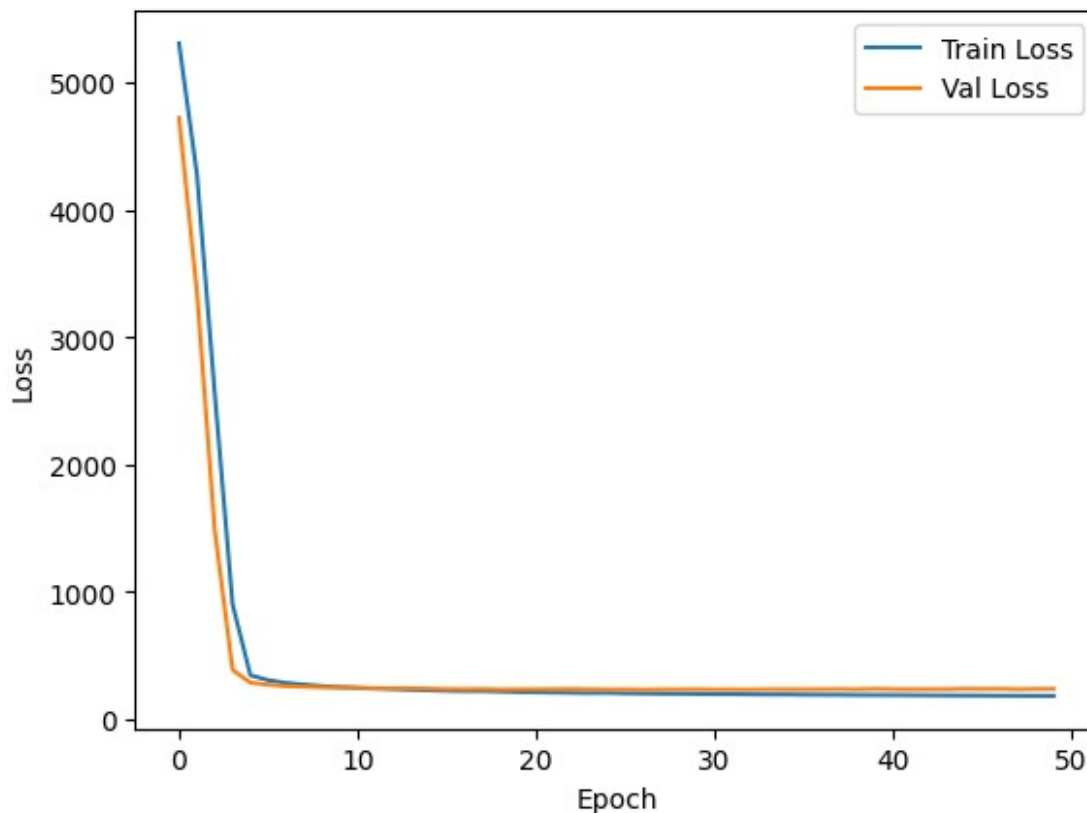


```
y_pred = model_1.predict(x_test_scaled)
```

```
4/4 ————— 0s 18ms/step
```

Evaluate

```
train_loss = history.history['loss']  
val_loss = history.history['val_loss']  
plt.plot(train_loss, label="Train Loss")  
plt.plot(val_loss, label='Val Loss')  
plt.xlabel("Epoch")  
plt.ylabel("Loss")  
plt.legend()  
plt.show()
```



The loss curve for the functional model shows that both training and validation loss drop significantly, just like in the sequential model. The smaller gap between the two suggests better generalization. However, the quick decrease in the plot indicates that the model might be too simple for the data, which could mean it's not capturing all the necessary details.

```
R2 = r2_score(y_test, y_pred)  
MAE = mean_absolute_error(y_test, y_pred)  
MSE = mean_squared_error(y_test, y_pred)
```

```

print("R2 Score=", R2)
print("Mean Absolute Error=", MAE)
print("Mean Squared Error=", MSE)

R2 Score= 0.10655987020032043
Mean Absolute Error= 13.075829557353057
Mean Squared Error= 316.3714876154902

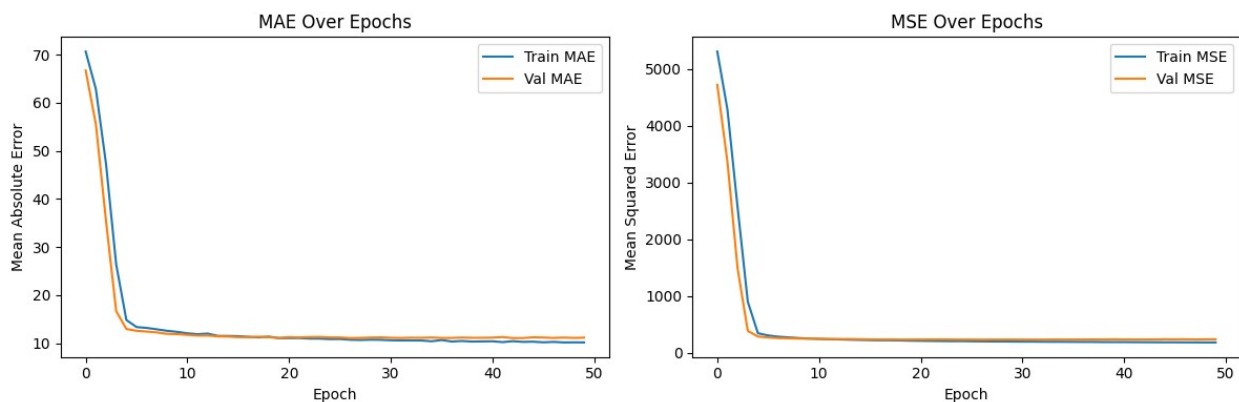
# Plot Mean Absolute Error (MAE)
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.plot(history.history['mean_absolute_error'], label='Train MAE')
plt.plot(history.history['val_mean_absolute_error'], label='Val MAE')
plt.xlabel('Epoch')
plt.ylabel('Mean Absolute Error')
plt.title('MAE Over Epochs')
plt.legend()

# Plot Mean Squared Error (MSE)
plt.subplot(1, 2, 2)
plt.plot(history.history['mean_squared_error'], label='Train MSE')
plt.plot(history.history['val_mean_squared_error'], label='Val MSE')
plt.xlabel('Epoch')
plt.ylabel('Mean Squared Error')
plt.title('MSE Over Epochs')
plt.legend()

plt.tight_layout()
plt.show()

```



The Functional API model also learned quickly as it shown by the plot of loss curve stable and not overfitting. but the statistical result is

1. $R^2 = 0.106$,
2. $MAE = 13.07$, and
3. $MSE = 316.37$

which means the model's predictive capability is low, with an R^2 score of 0.106, indicating that it explains only 10.6% of the variance in the data. The high values of MSE and MAE further indicate that the model struggles to provide accurate predictions, as it has a substantial average error of approximately 13.07 units.

Comparison Both Baseline Model

Both the Sequential and Functional API models displayed similar learning curves with a rapid reduction in loss, indicating that they both learned from the training data quickly. However, the low values of R^2 , high values of MAE, and MSE reveal that both models struggle to accurately predict productivity scores based on the data. Overall, the statistical results show that the Sequential model performs slightly better than the Functional API model, yielding more accurate predictions with lower error metrics.

In conclusion, both models appear to be too simplistic for the dataset, making it challenging to effectively predict productivity scores.

Modify Model

Next, I modified the ANN model to achieve better performance compared to the baseline model by adding Dropout for regularization to reduce overfitting, and BatchNormalization to stabilize and speed up the training process, resulting in more efficient learning and improved generalization.

Sequential

```
from tensorflow.keras.layers import Dense, Dropout, BatchNormalization
from tensorflow.keras.regularizers import l2

model_seq = Sequential([
    Dense(128, activation='elu', kernel_regularizer=l2(0.001),
input_shape=(22,)),
    BatchNormalization(),
    Dropout(0.3),

    Dense(64, activation='elu', kernel_regularizer=l2(0.001)),
    BatchNormalization(),
    Dropout(0.3),

    Dense(64, activation='elu', kernel_regularizer=l2(0.001)),
    BatchNormalization(),
    Dropout(0.3),

    Dense(64, activation='elu', kernel_regularizer=l2(0.001)),
    BatchNormalization(),

    Dense(1, activation='linear')
])
```

To enhance the performance of the sequential model, I made several modifications. I replaced the activation function with the Exponential Linear Unit (ELU) for each layer. ELU is advantageous because it helps mitigate the risk of vanishing gradients and avoids the "dying ReLU" problem, where neurons become inactive during training. Unlike ReLU, ELU allows for negative values, maintaining a mean output closer to zero, which supports better gradient flow and faster convergence.

Additionally, I implemented L2 regularization with a strength of 0.001 to help prevent overfitting by penalizing large weights. Batch normalization is retained after each dense layer to stabilize learning and improve convergence further. Dropout layers are included to randomly disable a portion of neurons during training, further reducing the likelihood of overfitting. After experimenting with both activation functions, I found that using ELU yielded better results, with lower loss and improved accuracy during training and validation. Consequently, this modified structure aims to enhance the overall performance of the model while effectively addressing potential overfitting issues.

```
model_seq.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),
    loss='mean_squared_error',
    metrics=['mean_absolute_error', 'mean_squared_error']
)
```

Also, using a learning rate of 0.001 for the optimizer helps to control the step size during training, allowing the model to converge more smoothly and efficiently.

Model Summary

```
model_seq.summary()
```

Model: "sequential_1"

Layer (type) Param #	Output Shape	
dense_6 (Dense) 2,944	(None, 128)	
batch_normalization (BatchNormalization) 512	(None, 128)	
dropout (Dropout) 0	(None, 128)	

dense_7 (Dense)	(None, 64)	
8,256		
batch_normalization_1	(None, 64)	
256		
(BatchNormalization)		
dropout_1 (Dropout)	(None, 64)	
0		
dense_8 (Dense)	(None, 64)	
4,160		
batch_normalization_2	(None, 64)	
256		
(BatchNormalization)		
dropout_2 (Dropout)	(None, 64)	
0		
dense_9 (Dense)	(None, 64)	
4,160		
batch_normalization_3	(None, 64)	
256		
(BatchNormalization)		
dense_10 (Dense)	(None, 1)	
65		

Total params: 20,865 (81.50 KB)

Trainable params: 20,225 (79.00 KB)

Non-trainable params: 640 (2.50 KB)

Then, I trained the model with EarlyStopping (patience = 10), which halts training if the validation loss doesn't improve for 10 consecutive epochs, and used 250 epochs to give the model enough time to converge while preventing overfitting.

```
early_stopping = EarlyStopping(monitor='val_loss', patience=10,  
restore_best_weights=True)
```

```
# Train the model
```

```
history = model_seq.fit(x_train_scaled, y_train,  
                        validation_data=(x_val_scaled, y_val),  
                        epochs=250,  
                        batch_size=32,  
                        verbose=1,  
                        callbacks=[early_stopping])
```

Epoch 1/250

```
26/26 ————— 5s 24ms/step - loss: 5724.2617 -  
mean_absolute_error: 73.7766 - mean_squared_error: 5724.0103 -  
val_loss: 5443.7822 - val_mean_absolute_error: 72.0558 -  
val_mean_squared_error: 5443.5308
```

Epoch 2/250

```
26/26 ————— 0s 8ms/step - loss: 5467.1777 -  
mean_absolute_error: 72.0414 - mean_squared_error: 5466.9258 -  
val_loss: 5380.5049 - val_mean_absolute_error: 71.7140 -  
val_mean_squared_error: 5380.2534
```

Epoch 3/250

```
26/26 ————— 0s 12ms/step - loss: 5511.7173 -  
mean_absolute_error: 72.4151 - mean_squared_error: 5511.4653 -  
val_loss: 5300.6914 - val_mean_absolute_error: 71.2235 -  
val_mean_squared_error: 5300.4380
```

Epoch 4/250

```
26/26 ————— 1s 15ms/step - loss: 5407.8628 -  
mean_absolute_error: 71.7485 - mean_squared_error: 5407.6099 -  
val_loss: 5169.8657 - val_mean_absolute_error: 70.3036 -  
val_mean_squared_error: 5169.6118
```

Epoch 5/250

```
26/26 ————— 1s 15ms/step - loss: 5281.1406 -  
mean_absolute_error: 70.6701 - mean_squared_error: 5280.8867 -  
val_loss: 5048.1245 - val_mean_absolute_error: 69.4427 -  
val_mean_squared_error: 5047.8701
```

Epoch 6/250

```
26/26 ————— 1s 12ms/step - loss: 5182.4531 -  
mean_absolute_error: 70.0996 - mean_squared_error: 5182.1987 -  
val_loss: 4861.6812 - val_mean_absolute_error: 68.0867 -  
val_mean_squared_error: 4861.4258
```

Epoch 7/250

```
26/26 ————— 1s 10ms/step - loss: 4943.2207 -  
mean_absolute_error: 68.4068 - mean_squared_error: 4942.9644 -  
val_loss: 4741.4385 - val_mean_absolute_error: 67.2018 -  
val_mean_squared_error: 4741.1826
```

Epoch 8/250

26/26 ————— 0s 8ms/step - loss: 4779.3091 -
mean_absolute_error: 67.3543 - mean_squared_error: 4779.0532 -
val_loss: 4507.3745 - val_mean_absolute_error: 65.4535 -
val_mean_squared_error: 4507.1172

Epoch 9/250

26/26 ————— 0s 10ms/step - loss: 4650.5112 -
mean_absolute_error: 66.3739 - mean_squared_error: 4650.2534 -
val_loss: 4309.6143 - val_mean_absolute_error: 63.9513 -
val_mean_squared_error: 4309.3555

Epoch 10/250

26/26 ————— 0s 8ms/step - loss: 4418.7729 -
mean_absolute_error: 64.6674 - mean_squared_error: 4418.5146 -
val_loss: 4041.0376 - val_mean_absolute_error: 61.8022 -
val_mean_squared_error: 4040.7786

Epoch 11/250

26/26 ————— 0s 10ms/step - loss: 4089.6418 -
mean_absolute_error: 62.0279 - mean_squared_error: 4089.3823 -
val_loss: 3767.1206 - val_mean_absolute_error: 59.5700 -
val_mean_squared_error: 3766.8606

Epoch 12/250

26/26 ————— 0s 10ms/step - loss: 4016.5298 -
mean_absolute_error: 61.5758 - mean_squared_error: 4016.2695 -
val_loss: 3499.3628 - val_mean_absolute_error: 57.2375 -
val_mean_squared_error: 3499.1018

Epoch 13/250

26/26 ————— 0s 8ms/step - loss: 3713.9653 -
mean_absolute_error: 58.9370 - mean_squared_error: 3713.7041 -
val_loss: 3276.9456 - val_mean_absolute_error: 55.3109 -
val_mean_squared_error: 3276.6833

Epoch 14/250

26/26 ————— 0s 8ms/step - loss: 3277.9785 -
mean_absolute_error: 55.2334 - mean_squared_error: 3277.7166 -
val_loss: 3000.5093 - val_mean_absolute_error: 52.6941 -
val_mean_squared_error: 3000.2458

Epoch 15/250

26/26 ————— 0s 7ms/step - loss: 3058.1572 -
mean_absolute_error: 53.2022 - mean_squared_error: 3057.8943 -
val_loss: 2714.4734 - val_mean_absolute_error: 49.9686 -
val_mean_squared_error: 2714.2092

Epoch 16/250

26/26 ————— 0s 7ms/step - loss: 2802.2651 -
mean_absolute_error: 50.6472 - mean_squared_error: 2802.0007 -
val_loss: 2460.8271 - val_mean_absolute_error: 47.4094 -
val_mean_squared_error: 2460.5623

Epoch 17/250

26/26 ————— 0s 8ms/step - loss: 2495.0017 -
mean_absolute_error: 47.6535 - mean_squared_error: 2494.7368 -
val_loss: 2157.6091 - val_mean_absolute_error: 44.1781 -

```
val_mean_squared_error: 2157.3438
Epoch 18/250
26/26 _____ 0s 7ms/step - loss: 2209.6882 -
mean_absolute_error: 44.5319 - mean_squared_error: 2209.4224 -
val_loss: 1907.0280 - val_mean_absolute_error: 41.3312 -
val_mean_squared_error: 1906.7616
Epoch 19/250
26/26 _____ 0s 8ms/step - loss: 2023.2583 -
mean_absolute_error: 42.4075 - mean_squared_error: 2022.9917 -
val_loss: 1697.6505 - val_mean_absolute_error: 38.8011 -
val_mean_squared_error: 1697.3833
Epoch 20/250
26/26 _____ 0s 8ms/step - loss: 1823.8800 -
mean_absolute_error: 40.0505 - mean_squared_error: 1823.6125 -
val_loss: 1451.3661 - val_mean_absolute_error: 35.6684 -
val_mean_squared_error: 1451.0978
Epoch 21/250
26/26 _____ 0s 7ms/step - loss: 1524.1216 -
mean_absolute_error: 36.3759 - mean_squared_error: 1523.8533 -
val_loss: 1262.1989 - val_mean_absolute_error: 33.0981 -
val_mean_squared_error: 1261.9298
Epoch 22/250
26/26 _____ 0s 8ms/step - loss: 1326.7906 -
mean_absolute_error: 33.6771 - mean_squared_error: 1326.5215 -
val_loss: 1083.7455 - val_mean_absolute_error: 30.3992 -
val_mean_squared_error: 1083.4756
Epoch 23/250
26/26 _____ 0s 9ms/step - loss: 1130.7609 -
mean_absolute_error: 31.0235 - mean_squared_error: 1130.4908 -
val_loss: 951.7398 - val_mean_absolute_error: 28.4463 -
val_mean_squared_error: 951.4692
Epoch 24/250
26/26 _____ 0s 9ms/step - loss: 985.9446 -
mean_absolute_error: 28.7602 - mean_squared_error: 985.6738 -
val_loss: 802.2562 - val_mean_absolute_error: 26.0437 -
val_mean_squared_error: 801.9848
Epoch 25/250
26/26 _____ 0s 9ms/step - loss: 871.2469 -
mean_absolute_error: 26.7054 - mean_squared_error: 870.9752 -
val_loss: 684.6730 - val_mean_absolute_error: 23.9663 -
val_mean_squared_error: 684.4006
Epoch 26/250
26/26 _____ 0s 9ms/step - loss: 720.0665 -
mean_absolute_error: 24.0708 - mean_squared_error: 719.7939 -
val_loss: 585.7772 - val_mean_absolute_error: 21.8152 -
val_mean_squared_error: 585.5040
Epoch 27/250
26/26 _____ 0s 9ms/step - loss: 574.5822 -
mean_absolute_error: 21.2405 - mean_squared_error: 574.3087 -
```

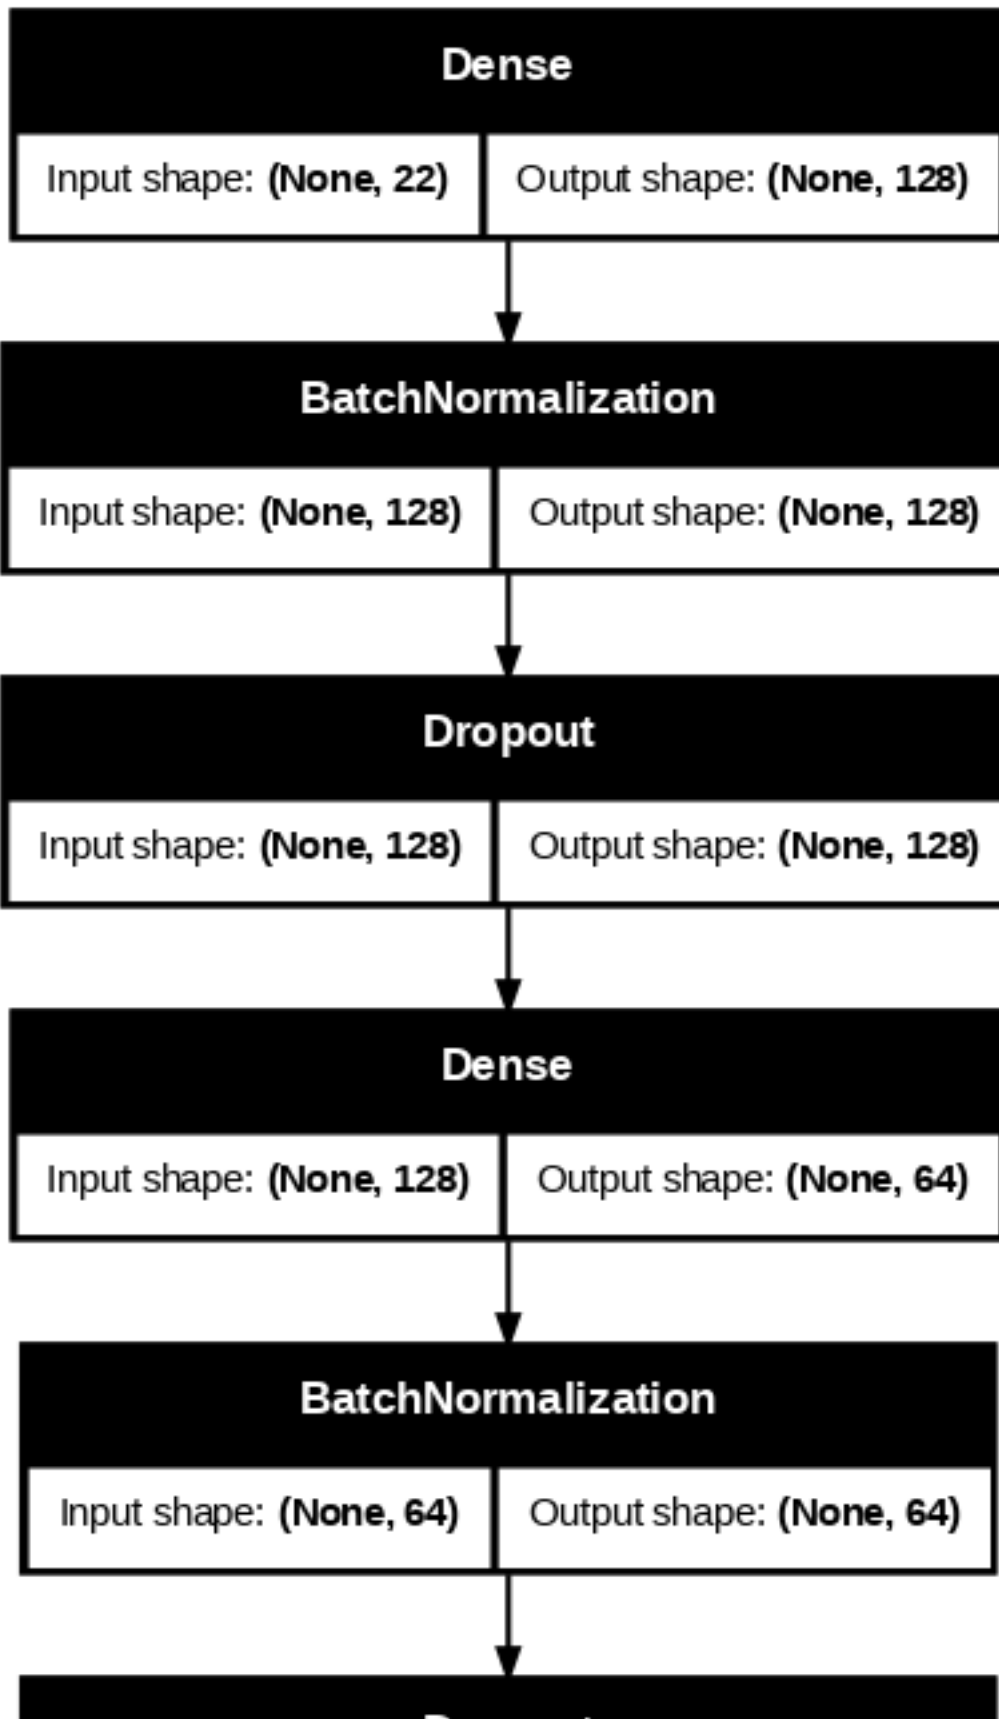
```
val_loss: 500.5959 - val_mean_absolute_error: 19.9472 -  
val_mean_squared_error: 500.3216  
Epoch 28/250  
26/26 ━━━━━━━━━━━ 0s 8ms/step - loss: 534.4766 -  
mean_absolute_error: 20.4113 - mean_squared_error: 534.2021 -  
val_loss: 428.6069 - val_mean_absolute_error: 18.3649 -  
val_mean_squared_error: 428.3318  
Epoch 29/250  
26/26 ━━━━━━━━━━━ 0s 9ms/step - loss: 459.3269 -  
mean_absolute_error: 18.6241 - mean_squared_error: 459.0515 -  
val_loss: 384.7834 - val_mean_absolute_error: 17.1221 -  
val_mean_squared_error: 384.5074  
Epoch 30/250  
26/26 ━━━━━━━━━━━ 0s 9ms/step - loss: 407.4894 -  
mean_absolute_error: 17.4938 - mean_squared_error: 407.2133 -  
val_loss: 350.3765 - val_mean_absolute_error: 16.1492 -  
val_mean_squared_error: 350.0997  
Epoch 31/250  
26/26 ━━━━━━━━━━━ 0s 9ms/step - loss: 335.8443 -  
mean_absolute_error: 15.5015 - mean_squared_error: 335.5672 -  
val_loss: 315.8338 - val_mean_absolute_error: 15.0055 -  
val_mean_squared_error: 315.5560  
Epoch 32/250  
26/26 ━━━━━━━━━━━ 0s 10ms/step - loss: 334.3896 -  
mean_absolute_error: 15.5895 - mean_squared_error: 334.1115 -  
val_loss: 280.6640 - val_mean_absolute_error: 13.7522 -  
val_mean_squared_error: 280.3853  
Epoch 33/250  
26/26 ━━━━━━━━━━━ 0s 9ms/step - loss: 289.1167 -  
mean_absolute_error: 14.2967 - mean_squared_error: 288.8378 -  
val_loss: 258.1236 - val_mean_absolute_error: 12.9254 -  
val_mean_squared_error: 257.8441  
Epoch 34/250  
26/26 ━━━━━━━━━━━ 0s 9ms/step - loss: 280.9151 -  
mean_absolute_error: 14.0128 - mean_squared_error: 280.6354 -  
val_loss: 243.9005 - val_mean_absolute_error: 12.5152 -  
val_mean_squared_error: 243.6201  
Epoch 35/250  
26/26 ━━━━━━━━━━━ 0s 9ms/step - loss: 236.2144 -  
mean_absolute_error: 12.4882 - mean_squared_error: 235.9337 -  
val_loss: 225.5256 - val_mean_absolute_error: 11.5096 -  
val_mean_squared_error: 225.2443  
Epoch 36/250  
26/26 ━━━━━━━━━━━ 0s 9ms/step - loss: 236.3746 -  
mean_absolute_error: 12.4649 - mean_squared_error: 236.0932 -  
val_loss: 219.6480 - val_mean_absolute_error: 11.3971 -  
val_mean_squared_error: 219.3660  
Epoch 37/250  
26/26 ━━━━━━━━━━━ 0s 9ms/step - loss: 238.5324 -
```

```
mean_absolute_error: 12.3212 - mean_squared_error: 238.2501 -  
val_loss: 215.4549 - val_mean_absolute_error: 11.1420 -  
val_mean_squared_error: 215.1719  
Epoch 38/250  
26/26 ————— 0s 8ms/step - loss: 227.9316 -  
mean_absolute_error: 11.9792 - mean_squared_error: 227.6484 -  
val_loss: 211.9232 - val_mean_absolute_error: 10.9186 -  
val_mean_squared_error: 211.6394  
Epoch 39/250  
26/26 ————— 0s 10ms/step - loss: 236.2992 -  
mean_absolute_error: 11.9718 - mean_squared_error: 236.0152 -  
val_loss: 208.4695 - val_mean_absolute_error: 10.6013 -  
val_mean_squared_error: 208.1848  
Epoch 40/250  
26/26 ————— 0s 9ms/step - loss: 211.3193 -  
mean_absolute_error: 11.2907 - mean_squared_error: 211.0345 -  
val_loss: 205.8438 - val_mean_absolute_error: 10.4031 -  
val_mean_squared_error: 205.5585  
Epoch 41/250  
26/26 ————— 0s 10ms/step - loss: 206.9157 -  
mean_absolute_error: 10.8756 - mean_squared_error: 206.6301 -  
val_loss: 202.3220 - val_mean_absolute_error: 10.1489 -  
val_mean_squared_error: 202.0357  
Epoch 42/250  
26/26 ————— 0s 11ms/step - loss: 188.4791 -  
mean_absolute_error: 10.8830 - mean_squared_error: 188.1927 -  
val_loss: 209.5815 - val_mean_absolute_error: 10.5258 -  
val_mean_squared_error: 209.2945  
Epoch 43/250  
26/26 ————— 0s 11ms/step - loss: 205.8478 -  
mean_absolute_error: 10.7959 - mean_squared_error: 205.5607 -  
val_loss: 201.3459 - val_mean_absolute_error: 10.0688 -  
val_mean_squared_error: 201.0582  
Epoch 44/250  
26/26 ————— 1s 11ms/step - loss: 213.5929 -  
mean_absolute_error: 11.1385 - mean_squared_error: 213.3049 -  
val_loss: 203.2230 - val_mean_absolute_error: 10.1354 -  
val_mean_squared_error: 202.9344  
Epoch 45/250  
26/26 ————— 0s 12ms/step - loss: 210.2280 -  
mean_absolute_error: 11.0206 - mean_squared_error: 209.9391 -  
val_loss: 205.6993 - val_mean_absolute_error: 10.1175 -  
val_mean_squared_error: 205.4099  
Epoch 46/250  
26/26 ————— 0s 15ms/step - loss: 198.6393 -  
mean_absolute_error: 10.5259 - mean_squared_error: 198.3496 -  
val_loss: 204.2334 - val_mean_absolute_error: 10.1237 -  
val_mean_squared_error: 203.9433  
Epoch 47/250
```


26/26 ————— 0s 10ms/step - loss: 186.1339 -
mean_absolute_error: 10.3697 - mean_squared_error: 185.8436 -
val_loss: 200.2683 - val_mean_absolute_error: 9.9804 -
val_mean_squared_error: 199.9776
Epoch 48/250
26/26 ————— 0s 9ms/step - loss: 216.7175 -
mean_absolute_error: 11.0761 - mean_squared_error: 216.4266 -
val_loss: 199.4671 - val_mean_absolute_error: 10.0266 -
val_mean_squared_error: 199.1758
Epoch 49/250
26/26 ————— 0s 9ms/step - loss: 187.4076 -
mean_absolute_error: 10.3623 - mean_squared_error: 187.1161 -
val_loss: 202.3458 - val_mean_absolute_error: 9.8667 -
val_mean_squared_error: 202.0537
Epoch 50/250
26/26 ————— 0s 8ms/step - loss: 195.7316 -
mean_absolute_error: 10.4226 - mean_squared_error: 195.4393 -
val_loss: 204.8990 - val_mean_absolute_error: 9.8820 -
val_mean_squared_error: 204.6062
Epoch 51/250
26/26 ————— 0s 7ms/step - loss: 182.9145 -
mean_absolute_error: 10.3950 - mean_squared_error: 182.6216 -
val_loss: 207.4431 - val_mean_absolute_error: 10.0829 -
val_mean_squared_error: 207.1497
Epoch 52/250
26/26 ————— 0s 8ms/step - loss: 179.6785 -
mean_absolute_error: 10.2292 - mean_squared_error: 179.3851 -
val_loss: 208.4056 - val_mean_absolute_error: 10.0830 -
val_mean_squared_error: 208.1118
Epoch 53/250
26/26 ————— 0s 7ms/step - loss: 199.2600 -
mean_absolute_error: 10.5399 - mean_squared_error: 198.9659 -
val_loss: 204.7027 - val_mean_absolute_error: 9.9881 -
val_mean_squared_error: 204.4080
Epoch 54/250
26/26 ————— 0s 8ms/step - loss: 185.1533 -
mean_absolute_error: 10.2468 - mean_squared_error: 184.8585 -
val_loss: 205.8557 - val_mean_absolute_error: 10.0166 -
val_mean_squared_error: 205.5604
Epoch 55/250
26/26 ————— 0s 7ms/step - loss: 175.9093 -
mean_absolute_error: 10.0490 - mean_squared_error: 175.6138 -
val_loss: 204.2377 - val_mean_absolute_error: 9.9379 -
val_mean_squared_error: 203.9418
Epoch 56/250
26/26 ————— 0s 8ms/step - loss: 188.7411 -
mean_absolute_error: 10.1233 - mean_squared_error: 188.4451 -
val_loss: 202.5629 - val_mean_absolute_error: 9.9348 -
val_mean_squared_error: 202.2664

```
Epoch 57/250
26/26 _____ 0s 9ms/step - loss: 181.0585 -
mean_absolute_error: 10.1216 - mean_squared_error: 180.7618 -
val_loss: 203.5051 - val_mean_absolute_error: 9.9108 -
val_mean_squared_error: 203.2078
Epoch 58/250
26/26 _____ 0s 8ms/step - loss: 207.5447 -
mean_absolute_error: 10.9342 - mean_squared_error: 207.2473 -
val_loss: 205.9593 - val_mean_absolute_error: 9.9551 -
val_mean_squared_error: 205.6613

plot_model(model_seq, show_shapes=True, dpi=75)
```



```
# Prediction
```

```
y_pred = model_seq.predict(x_test_scaled)
```

WARNING:tensorflow:5 out of the last 9 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x7c8d618e7880> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

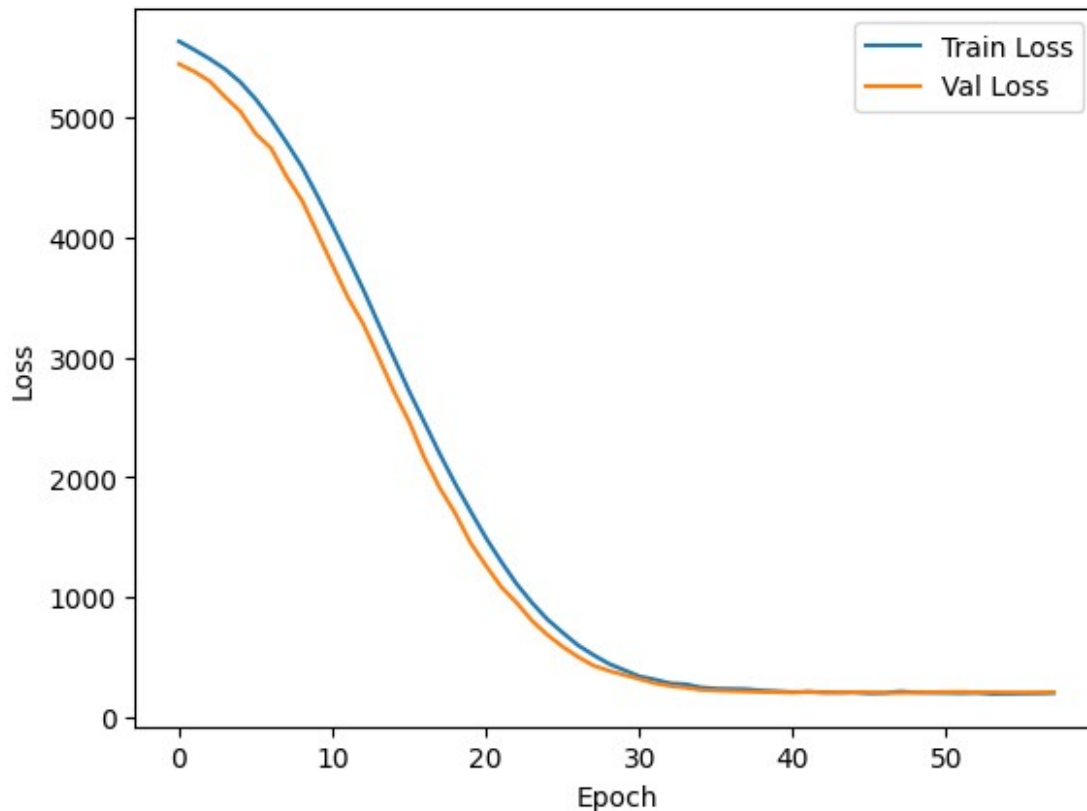
1/4 ————— 0s 149ms/step

WARNING:tensorflow:6 out of the last 12 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x7c8d618e7880> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

4/4 ————— 0s 54ms/step

Evaluate

```
train_loss = history.history['loss']
val_loss = history.history['val_loss']
plt.plot(train_loss, label="Train Loss")
plt.plot(val_loss, label='Val Loss')
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



This loss curve plot is much improved, as the model trains more effectively. The loss starts to decline at around epoch 30 and stabilizes between epochs 40 and 50, indicating that the model is converging. The closeness of the training and validation loss suggests that the model is generalizing well to the validation data without significant overfitting. Overall, this behavior reflects that the adjustments made to the model such as using ELU activation and incorporating L2 regularization are having a positive impact on training dynamics and model performance.

```
R2 = r2_score(y_test, y_pred)
MAE = mean_absolute_error(y_test, y_pred)
MSE = mean_squared_error(y_test, y_pred)

print("R² Score=", R2)
print("Mean Absolute Error=", MAE)
print("Mean Squared Error=", MSE)

R² Score= 0.25563211127907237
Mean Absolute Error= 11.70159676440009
Mean Squared Error= 263.58428330350785

# Plot Mean Absolute Error (MAE)
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.plot(history.history['mean_absolute_error'], label='Train MAE')
plt.plot(history.history['val_mean_absolute_error'], label='Val MAE')
```

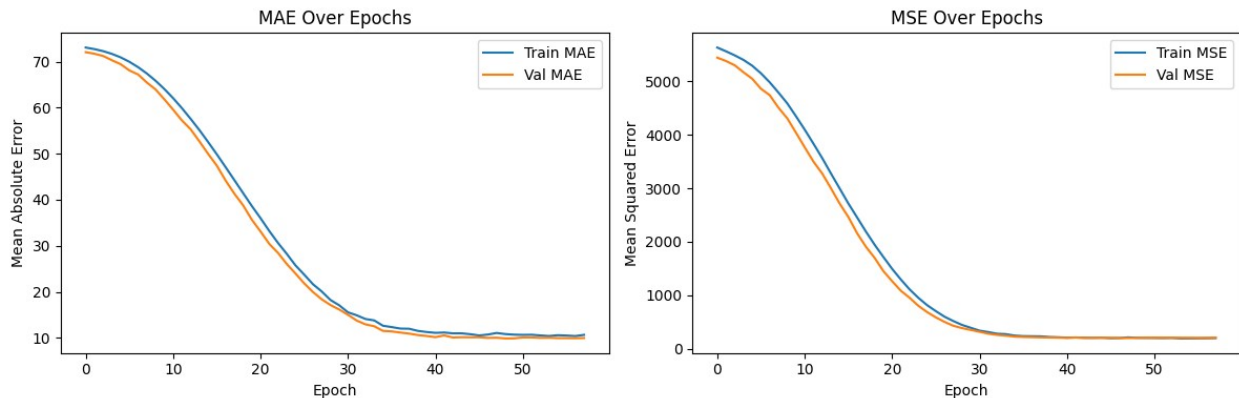
```

plt.xlabel('Epoch')
plt.ylabel('Mean Absolute Error')
plt.title('MAE Over Epochs')
plt.legend()

# Plot Mean Squared Error (MSE)
plt.subplot(1, 2, 2)
plt.plot(history.history['mean_squared_error'], label='Train MSE')
plt.plot(history.history['val_mean_squared_error'], label='Val MSE')
plt.xlabel('Epoch')
plt.ylabel('Mean Squared Error')
plt.title('MSE Over Epochs')
plt.legend()

plt.tight_layout()
plt.show()

```



The modified Sequential model shows a relatively low R^2 score of 0.25, meaning it only explains 25% of the variance in the target variable. Additionally, the high MAE (11.7) and MSE (263.58) indicate that the model's predictions are not very close to the actual values. A low R^2 combined with high error values suggests that the model struggles to capture the underlying patterns in the data and may not be complex enough or is missing important features needed for accurate prediction.

Functional API

```

from tensorflow.keras.layers import Input

inputs = Input(shape=(x_train_scaled.shape[1],))

# Input layer
inputs = Input(shape=(22,))

# Hidden layers
x = Dense(128, activation='elu', kernel_regularizer=l2(0.001))(inputs)
x = BatchNormalization()(x)
x = Dropout(0.3)(x)

```

```

x = Dense(64, activation='elu', kernel_regularizer=l2(0.001))(x)
x = BatchNormalization()(x)
x = Dropout(0.3)(x)

x = Dense(64, activation='elu', kernel_regularizer=l2(0.001))(x)
x = BatchNormalization()(x)
x = Dropout(0.3)(x)

x = Dense(64, activation='elu', kernel_regularizer=l2(0.001))(x)
x = BatchNormalization()(x)

# Output layer
outputs = Dense(1, activation='linear')(x)

# Create model
model_func = Model(inputs=inputs, outputs=outputs)

```

This uses a Functional API model with a series of Dense layers and ReLU activation functions in the hidden layers to learn non-linear patterns. The model starts with an input layer that accepts 25 features. Each hidden layer uses BatchNormalization to stabilize training and Dropout (0.3 rate) to prevent overfitting. The architecture has four hidden layers with 128, 64, 64, and 64 units, respectively, and concludes with a linear activation function in the output layer, making it suitable for regression tasks.

```

model_func.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),
    loss='mean_squared_error',
    metrics=['mean_absolute_error', 'mean_squared_error']
)

```

Additionally, the model is compiled with an Adam optimizer using a learning rate of 0.001, which helps to control the step size during training, ensuring efficient convergence.

```

# Show model summary
model_func.summary()

```

Model: "functional_3"

Layer (type)	Output Shape
Param #	
input_layer_4 (InputLayer)	(None, 22)
0	

dense_11 (Dense)	(None, 128)	
2,944		
batch_normalization_4	(None, 128)	
512		
(BatchNormalization)		
dropout_3 (Dropout)	(None, 128)	
0		
dense_12 (Dense)	(None, 64)	
8,256		
batch_normalization_5	(None, 64)	
256		
(BatchNormalization)		
dropout_4 (Dropout)	(None, 64)	
0		
dense_13 (Dense)	(None, 64)	
4,160		
batch_normalization_6	(None, 64)	
256		
(BatchNormalization)		
dropout_5 (Dropout)	(None, 64)	
0		
dense_14 (Dense)	(None, 64)	
4,160		
batch_normalization_7	(None, 64)	
256		
(BatchNormalization)		

dense_15 (Dense)	(None, 1)	
65		

Total params: 20,865 (81.50 KB)

Trainable params: 20,225 (79.00 KB)

Non-trainable params: 640 (2.50 KB)

```
early_stopping = EarlyStopping(monitor='val_loss', patience=10,
                                restore_best_weights=True)
```

Train the model

```
history = model_func.fit(x_train_scaled, y_train,
                        validation_data=(x_val_scaled, y_val),
                        epochs=250,
                        batch_size=32,
                        verbose=1,
                        callbacks=[early_stopping])
```

Epoch 1/250

```
26/26 ————— 6s 34ms/step - loss: 5665.9287 -
mean_absolute_error: 73.1542 - mean_squared_error: 5665.6777 -
val_loss: 5429.6685 - val_mean_absolute_error: 71.9804 -
val_mean_squared_error: 5429.4170
```

Epoch 2/250

```
26/26 ————— 1s 12ms/step - loss: 5512.6455 -
mean_absolute_error: 72.4187 - mean_squared_error: 5512.3931 -
val_loss: 5365.8384 - val_mean_absolute_error: 71.6263 -
val_mean_squared_error: 5365.5859
```

Epoch 3/250

```
26/26 ————— 0s 13ms/step - loss: 5658.9316 -
mean_absolute_error: 73.4071 - mean_squared_error: 5658.6797 -
val_loss: 5251.0713 - val_mean_absolute_error: 70.8818 -
val_mean_squared_error: 5250.8184
```

Epoch 4/250

```
26/26 ————— 0s 12ms/step - loss: 5446.0249 -
mean_absolute_error: 71.9690 - mean_squared_error: 5445.7710 -
val_loss: 5132.5806 - val_mean_absolute_error: 70.0485 -
val_mean_squared_error: 5132.3271
```

Epoch 5/250

```
26/26 ————— 0s 7ms/step - loss: 5371.6499 -
mean_absolute_error: 71.5135 - mean_squared_error: 5371.3965 -
val_loss: 5047.9810 - val_mean_absolute_error: 69.4352 -
val_mean_squared_error: 5047.7261
```

Epoch 6/250

26/26 ————— 0s 8ms/step - loss: 5356.7476 -
mean_absolute_error: 71.4196 - mean_squared_error: 5356.4927 -
val_loss: 4864.0693 - val_mean_absolute_error: 68.1199 -
val_mean_squared_error: 4863.8145

Epoch 7/250

26/26 ————— 0s 9ms/step - loss: 5109.2930 -
mean_absolute_error: 69.8043 - mean_squared_error: 5109.0376 -
val_loss: 4621.4082 - val_mean_absolute_error: 66.3593 -
val_mean_squared_error: 4621.1519

Epoch 8/250

26/26 ————— 0s 9ms/step - loss: 4918.2568 -
mean_absolute_error: 68.3500 - mean_squared_error: 4917.9995 -
val_loss: 4402.0698 - val_mean_absolute_error: 64.6534 -
val_mean_squared_error: 4401.8120

Epoch 9/250

26/26 ————— 0s 9ms/step - loss: 4685.5166 -
mean_absolute_error: 66.6817 - mean_squared_error: 4685.2588 -
val_loss: 4210.3975 - val_mean_absolute_error: 63.1441 -
val_mean_squared_error: 4210.1382

Epoch 10/250

26/26 ————— 0s 9ms/step - loss: 4490.5869 -
mean_absolute_error: 65.3843 - mean_squared_error: 4490.3281 -
val_loss: 3941.9592 - val_mean_absolute_error: 60.9800 -
val_mean_squared_error: 3941.6997

Epoch 11/250

26/26 ————— 0s 9ms/step - loss: 4170.7676 -
mean_absolute_error: 62.6504 - mean_squared_error: 4170.5078 -
val_loss: 3734.9829 - val_mean_absolute_error: 59.2835 -
val_mean_squared_error: 3734.7219

Epoch 12/250

26/26 ————— 0s 9ms/step - loss: 4074.5413 -
mean_absolute_error: 61.9526 - mean_squared_error: 4074.2805 -
val_loss: 3481.3250 - val_mean_absolute_error: 57.0951 -
val_mean_squared_error: 3481.0630

Epoch 13/250

26/26 ————— 0s 9ms/step - loss: 3563.0049 -
mean_absolute_error: 57.6526 - mean_squared_error: 3562.7432 -
val_loss: 3215.4346 - val_mean_absolute_error: 54.7024 -
val_mean_squared_error: 3215.1714

Epoch 14/250

26/26 ————— 0s 10ms/step - loss: 3412.1855 -
mean_absolute_error: 56.4578 - mean_squared_error: 3411.9224 -
val_loss: 2999.7119 - val_mean_absolute_error: 52.7570 -
val_mean_squared_error: 2999.4480

Epoch 15/250

26/26 ————— 0s 10ms/step - loss: 3081.7068 -
mean_absolute_error: 53.4666 - mean_squared_error: 3081.4429 -
val_loss: 2697.0515 - val_mean_absolute_error: 49.8268 -

```
val_mean_squared_error: 2696.7869
Epoch 16/250
26/26 _____ 0s 9ms/step - loss: 2879.0903 -
mean_absolute_error: 51.4180 - mean_squared_error: 2878.8254 -
val_loss: 2438.2827 - val_mean_absolute_error: 47.2296 -
val_mean_squared_error: 2438.0168
Epoch 17/250
26/26 _____ 0s 9ms/step - loss: 2582.2778 -
mean_absolute_error: 48.4954 - mean_squared_error: 2582.0120 -
val_loss: 2157.3513 - val_mean_absolute_error: 44.1326 -
val_mean_squared_error: 2157.0845
Epoch 18/250
26/26 _____ 0s 10ms/step - loss: 2327.5178 -
mean_absolute_error: 45.5622 - mean_squared_error: 2327.2512 -
val_loss: 1932.2089 - val_mean_absolute_error: 41.6023 -
val_mean_squared_error: 1931.9413
Epoch 19/250
26/26 _____ 0s 9ms/step - loss: 1977.8993 -
mean_absolute_error: 41.7408 - mean_squared_error: 1977.6316 -
val_loss: 1670.6161 - val_mean_absolute_error: 38.4532 -
val_mean_squared_error: 1670.3477
Epoch 20/250
26/26 _____ 0s 10ms/step - loss: 1816.7772 -
mean_absolute_error: 40.1720 - mean_squared_error: 1816.5084 -
val_loss: 1444.9323 - val_mean_absolute_error: 35.5606 -
val_mean_squared_error: 1444.6626
Epoch 21/250
26/26 _____ 0s 8ms/step - loss: 1548.9656 -
mean_absolute_error: 36.4852 - mean_squared_error: 1548.6959 -
val_loss: 1273.0443 - val_mean_absolute_error: 33.3103 -
val_mean_squared_error: 1272.7738
Epoch 22/250
26/26 _____ 0s 8ms/step - loss: 1349.4875 -
mean_absolute_error: 33.7933 - mean_squared_error: 1349.2168 -
val_loss: 1106.5234 - val_mean_absolute_error: 30.8554 -
val_mean_squared_error: 1106.2518
Epoch 23/250
26/26 _____ 0s 7ms/step - loss: 1212.7385 -
mean_absolute_error: 32.0659 - mean_squared_error: 1212.4667 -
val_loss: 944.6947 - val_mean_absolute_error: 28.3365 -
val_mean_squared_error: 944.4224
Epoch 24/250
26/26 _____ 0s 7ms/step - loss: 1006.6848 -
mean_absolute_error: 29.0044 - mean_squared_error: 1006.4122 -
val_loss: 824.4277 - val_mean_absolute_error: 26.4052 -
val_mean_squared_error: 824.1545
Epoch 25/250
26/26 _____ 0s 8ms/step - loss: 877.6826 -
mean_absolute_error: 26.9926 - mean_squared_error: 877.4092 -
```

val_loss: 692.9370 - val_mean_absolute_error: 23.9764 -
val_mean_squared_error: 692.6630
Epoch 26/250
26/26 ————— 0s 8ms/step - loss: 737.8717 -
mean_absolute_error: 24.3291 - mean_squared_error: 737.5974 -
val_loss: 600.7045 - val_mean_absolute_error: 22.2789 -
val_mean_squared_error: 600.4296
Epoch 27/250
26/26 ————— 0s 7ms/step - loss: 613.8931 -
mean_absolute_error: 21.9988 - mean_squared_error: 613.6178 -
val_loss: 508.0497 - val_mean_absolute_error: 20.1819 -
val_mean_squared_error: 507.7736
Epoch 28/250
26/26 ————— 0s 7ms/step - loss: 555.0670 -
mean_absolute_error: 20.8545 - mean_squared_error: 554.7906 -
val_loss: 433.8089 - val_mean_absolute_error: 18.4337 -
val_mean_squared_error: 433.5317
Epoch 29/250
26/26 ————— 0s 9ms/step - loss: 476.7657 -
mean_absolute_error: 19.4838 - mean_squared_error: 476.4883 -
val_loss: 379.9834 - val_mean_absolute_error: 16.9539 -
val_mean_squared_error: 379.7053
Epoch 30/250
26/26 ————— 0s 10ms/step - loss: 394.8156 -
mean_absolute_error: 17.1929 - mean_squared_error: 394.5372 -
val_loss: 329.7207 - val_mean_absolute_error: 15.6565 -
val_mean_squared_error: 329.4415
Epoch 31/250
26/26 ————— 0s 8ms/step - loss: 362.0969 -
mean_absolute_error: 16.3908 - mean_squared_error: 361.8175 -
val_loss: 291.2494 - val_mean_absolute_error: 14.5457 -
val_mean_squared_error: 290.9693
Epoch 32/250
26/26 ————— 0s 8ms/step - loss: 322.6794 -
mean_absolute_error: 15.2596 - mean_squared_error: 322.3991 -
val_loss: 258.1195 - val_mean_absolute_error: 13.1898 -
val_mean_squared_error: 257.8384
Epoch 33/250
26/26 ————— 0s 7ms/step - loss: 284.5695 -
mean_absolute_error: 14.1321 - mean_squared_error: 284.2882 -
val_loss: 242.6507 - val_mean_absolute_error: 12.5396 -
val_mean_squared_error: 242.3686
Epoch 34/250
26/26 ————— 0s 8ms/step - loss: 261.2295 -
mean_absolute_error: 13.3538 - mean_squared_error: 260.9471 -
val_loss: 230.9603 - val_mean_absolute_error: 11.8417 -
val_mean_squared_error: 230.6773
Epoch 35/250
26/26 ————— 0s 7ms/step - loss: 245.7377 -

```
mean_absolute_error: 12.8542 - mean_squared_error: 245.4545 -  
val_loss: 223.9962 - val_mean_absolute_error: 11.4067 -  
val_mean_squared_error: 223.7124  
Epoch 36/250  
26/26 ————— 0s 7ms/step - loss: 260.6072 -  
mean_absolute_error: 13.1069 - mean_squared_error: 260.3231 -  
val_loss: 221.0378 - val_mean_absolute_error: 11.2504 -  
val_mean_squared_error: 220.7527  
Epoch 37/250  
26/26 ————— 0s 7ms/step - loss: 236.9454 -  
mean_absolute_error: 12.6195 - mean_squared_error: 236.6602 -  
val_loss: 215.0209 - val_mean_absolute_error: 10.7067 -  
val_mean_squared_error: 214.7351  
Epoch 38/250  
26/26 ————— 0s 8ms/step - loss: 214.9520 -  
mean_absolute_error: 11.4784 - mean_squared_error: 214.6659 -  
val_loss: 211.4196 - val_mean_absolute_error: 10.6173 -  
val_mean_squared_error: 211.1328  
Epoch 39/250  
26/26 ————— 0s 7ms/step - loss: 222.0407 -  
mean_absolute_error: 11.7771 - mean_squared_error: 221.7537 -  
val_loss: 206.2873 - val_mean_absolute_error: 10.3380 -  
val_mean_squared_error: 205.9998  
Epoch 40/250  
26/26 ————— 0s 8ms/step - loss: 225.8286 -  
mean_absolute_error: 11.6632 - mean_squared_error: 225.5409 -  
val_loss: 205.2498 - val_mean_absolute_error: 10.1172 -  
val_mean_squared_error: 204.9616  
Epoch 41/250  
26/26 ————— 0s 13ms/step - loss: 212.1503 -  
mean_absolute_error: 11.4280 - mean_squared_error: 211.8618 -  
val_loss: 207.1917 - val_mean_absolute_error: 10.1474 -  
val_mean_squared_error: 206.9026  
Epoch 42/250  
26/26 ————— 0s 13ms/step - loss: 197.7590 -  
mean_absolute_error: 10.5632 - mean_squared_error: 197.4697 -  
val_loss: 208.1175 - val_mean_absolute_error: 10.1674 -  
val_mean_squared_error: 207.8277  
Epoch 43/250  
26/26 ————— 1s 15ms/step - loss: 202.4885 -  
mean_absolute_error: 10.9343 - mean_squared_error: 202.1985 -  
val_loss: 208.3851 - val_mean_absolute_error: 10.1663 -  
val_mean_squared_error: 208.0944  
Epoch 44/250  
26/26 ————— 0s 16ms/step - loss: 221.7048 -  
mean_absolute_error: 11.4107 - mean_squared_error: 221.4139 -  
val_loss: 205.4678 - val_mean_absolute_error: 9.9670 -  
val_mean_squared_error: 205.1765  
Epoch 45/250
```

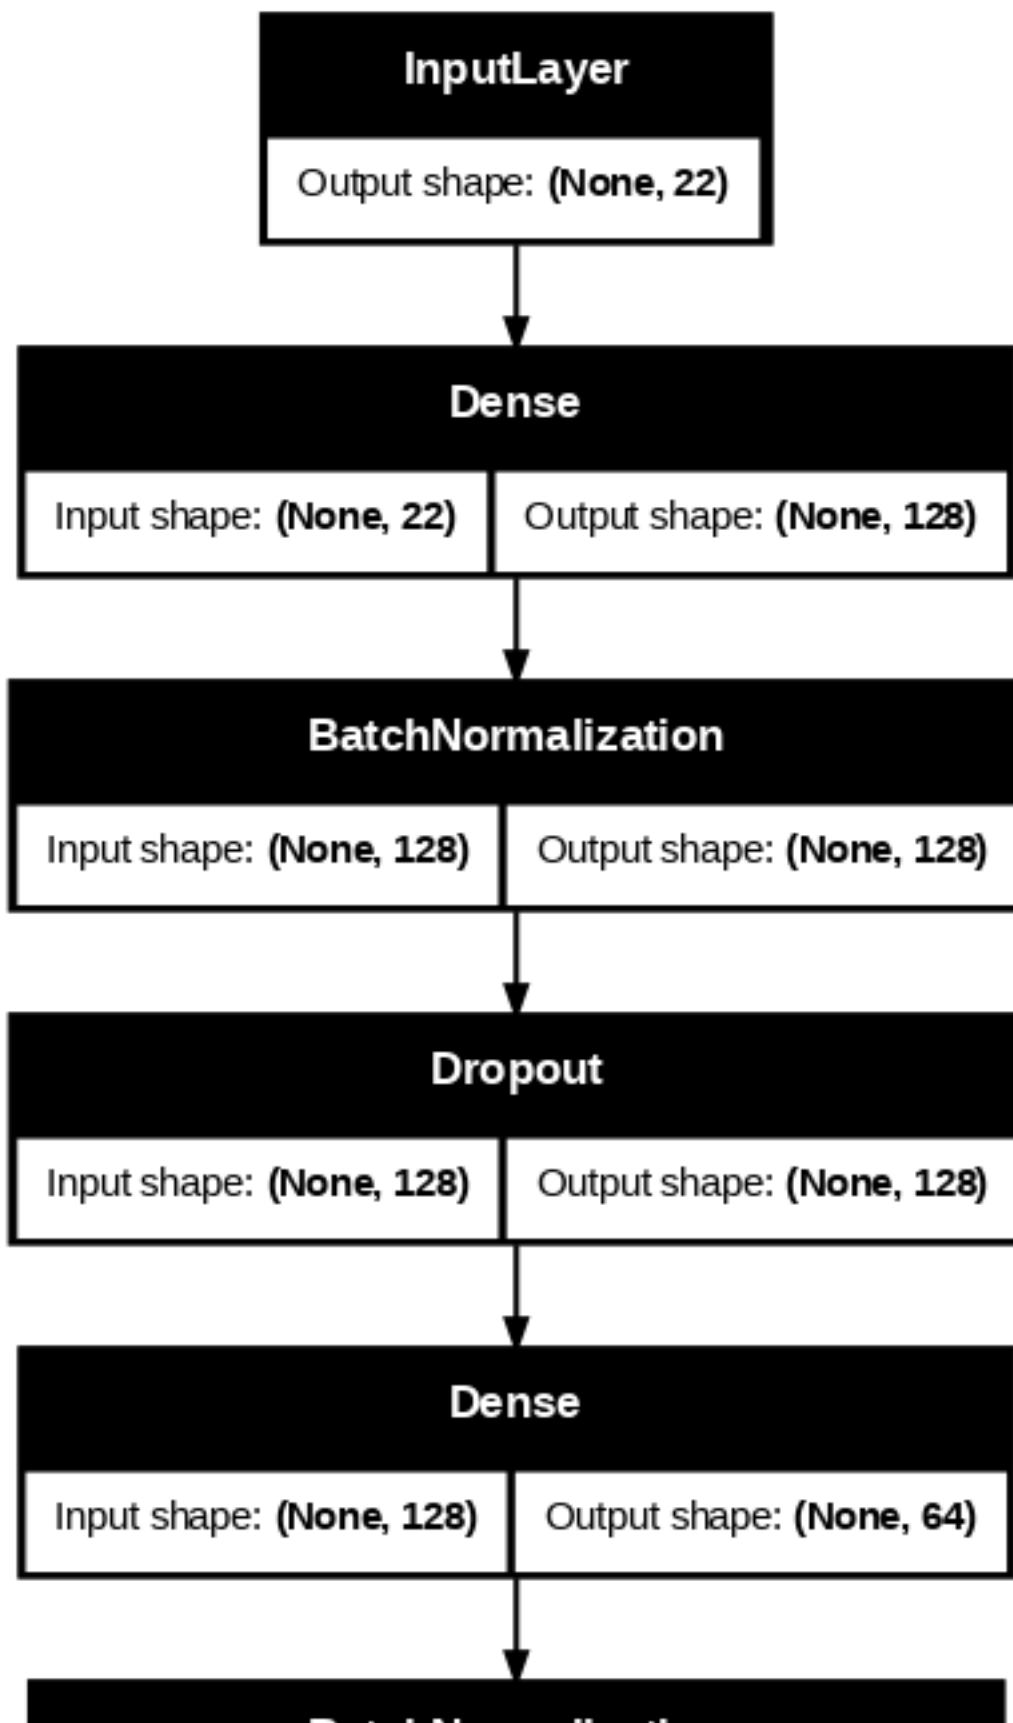
26/26 ————— 1s 19ms/step - loss: 207.1124 -
mean_absolute_error: 11.1623 - mean_squared_error: 206.8208 -
val_loss: 201.2608 - val_mean_absolute_error: 9.9199 -
val_mean_squared_error: 200.9688
Epoch 46/250
26/26 ————— 0s 11ms/step - loss: 193.7361 -
mean_absolute_error: 10.4989 - mean_squared_error: 193.4438 -
val_loss: 204.3364 - val_mean_absolute_error: 9.9699 -
val_mean_squared_error: 204.0435
Epoch 47/250
26/26 ————— 0s 7ms/step - loss: 201.8724 -
mean_absolute_error: 10.7469 - mean_squared_error: 201.5794 -
val_loss: 201.5531 - val_mean_absolute_error: 9.7041 -
val_mean_squared_error: 201.2596
Epoch 48/250
26/26 ————— 0s 8ms/step - loss: 187.1291 -
mean_absolute_error: 10.3637 - mean_squared_error: 186.8354 -
val_loss: 202.3413 - val_mean_absolute_error: 9.6968 -
val_mean_squared_error: 202.0470
Epoch 49/250
26/26 ————— 0s 8ms/step - loss: 214.0778 -
mean_absolute_error: 10.8448 - mean_squared_error: 213.7832 -
val_loss: 203.1346 - val_mean_absolute_error: 9.7665 -
val_mean_squared_error: 202.8394
Epoch 50/250
26/26 ————— 0s 7ms/step - loss: 198.3446 -
mean_absolute_error: 10.6867 - mean_squared_error: 198.0492 -
val_loss: 202.8846 - val_mean_absolute_error: 9.8146 -
val_mean_squared_error: 202.5884
Epoch 51/250
26/26 ————— 0s 7ms/step - loss: 188.2829 -
mean_absolute_error: 10.2407 - mean_squared_error: 187.9866 -
val_loss: 202.8741 - val_mean_absolute_error: 9.8426 -
val_mean_squared_error: 202.5772
Epoch 52/250
26/26 ————— 0s 8ms/step - loss: 199.5571 -
mean_absolute_error: 10.8990 - mean_squared_error: 199.2601 -
val_loss: 205.5434 - val_mean_absolute_error: 9.9500 -
val_mean_squared_error: 205.2458
Epoch 53/250
26/26 ————— 0s 8ms/step - loss: 196.6833 -
mean_absolute_error: 10.5712 - mean_squared_error: 196.3855 -
val_loss: 200.9202 - val_mean_absolute_error: 9.6747 -
val_mean_squared_error: 200.6218
Epoch 54/250
26/26 ————— 1s 26ms/step - loss: 205.8728 -
mean_absolute_error: 10.7653 - mean_squared_error: 205.5741 -
val_loss: 198.8845 - val_mean_absolute_error: 9.6965 -
val_mean_squared_error: 198.5854

Epoch 55/250
26/26 ————— 0s 7ms/step - loss: 209.3855 -
mean_absolute_error: 10.8626 - mean_squared_error: 209.0862 -
val_loss: 200.1332 - val_mean_absolute_error: 9.9024 -
val_mean_squared_error: 199.8335
Epoch 56/250
26/26 ————— 1s 22ms/step - loss: 193.5393 -
mean_absolute_error: 10.4671 - mean_squared_error: 193.2394 -
val_loss: 198.6743 - val_mean_absolute_error: 9.6895 -
val_mean_squared_error: 198.3739
Epoch 57/250
26/26 ————— 1s 24ms/step - loss: 188.1380 -
mean_absolute_error: 10.3194 - mean_squared_error: 187.8374 -
val_loss: 198.0440 - val_mean_absolute_error: 9.7075 -
val_mean_squared_error: 197.7428
Epoch 58/250
26/26 ————— 0s 7ms/step - loss: 201.6210 -
mean_absolute_error: 10.9581 - mean_squared_error: 201.3196 -
val_loss: 199.4951 - val_mean_absolute_error: 9.8770 -
val_mean_squared_error: 199.1931
Epoch 59/250
26/26 ————— 0s 9ms/step - loss: 201.9776 -
mean_absolute_error: 10.7121 - mean_squared_error: 201.6754 -
val_loss: 197.4078 - val_mean_absolute_error: 9.7034 -
val_mean_squared_error: 197.1051
Epoch 60/250
26/26 ————— 0s 10ms/step - loss: 199.7760 -
mean_absolute_error: 10.5606 - mean_squared_error: 199.4733 -
val_loss: 196.5782 - val_mean_absolute_error: 9.7409 -
val_mean_squared_error: 196.2749
Epoch 61/250
26/26 ————— 0s 9ms/step - loss: 190.6182 -
mean_absolute_error: 10.3645 - mean_squared_error: 190.3147 -
val_loss: 197.4910 - val_mean_absolute_error: 9.8230 -
val_mean_squared_error: 197.1870
Epoch 62/250
26/26 ————— 0s 9ms/step - loss: 194.5708 -
mean_absolute_error: 10.4087 - mean_squared_error: 194.2666 -
val_loss: 195.8092 - val_mean_absolute_error: 9.7470 -
val_mean_squared_error: 195.5045
Epoch 63/250
26/26 ————— 0s 8ms/step - loss: 186.0467 -
mean_absolute_error: 10.2564 - mean_squared_error: 185.7419 -
val_loss: 199.8962 - val_mean_absolute_error: 9.7496 -
val_mean_squared_error: 199.5908
Epoch 64/250
26/26 ————— 0s 8ms/step - loss: 183.3465 -
mean_absolute_error: 10.2090 - mean_squared_error: 183.0409 -
val_loss: 199.4720 - val_mean_absolute_error: 9.7988 -

```
val_mean_squared_error: 199.1659
Epoch 65/250
26/26 _____ 0s 7ms/step - loss: 177.3903 -
mean_absolute_error: 10.0301 - mean_squared_error: 177.0840 -
val_loss: 201.4003 - val_mean_absolute_error: 9.8751 -
val_mean_squared_error: 201.0936
Epoch 66/250
26/26 _____ 0s 9ms/step - loss: 177.5214 -
mean_absolute_error: 9.8270 - mean_squared_error: 177.2145 - val_loss:
199.1071 - val_mean_absolute_error: 9.7529 - val_mean_squared_error:
198.7996
Epoch 67/250
26/26 _____ 0s 9ms/step - loss: 188.4021 -
mean_absolute_error: 10.3005 - mean_squared_error: 188.0944 -
val_loss: 204.8605 - val_mean_absolute_error: 10.0210 -
val_mean_squared_error: 204.5523
Epoch 68/250
26/26 _____ 0s 8ms/step - loss: 195.3024 -
mean_absolute_error: 10.6146 - mean_squared_error: 194.9942 -
val_loss: 197.2351 - val_mean_absolute_error: 9.8085 -
val_mean_squared_error: 196.9263
Epoch 69/250
26/26 _____ 0s 8ms/step - loss: 181.7300 -
mean_absolute_error: 9.9151 - mean_squared_error: 181.4210 - val_loss:
195.6443 - val_mean_absolute_error: 9.7471 - val_mean_squared_error:
195.3348
Epoch 70/250
26/26 _____ 0s 7ms/step - loss: 192.4814 -
mean_absolute_error: 10.5670 - mean_squared_error: 192.1719 -
val_loss: 196.0975 - val_mean_absolute_error: 9.6810 -
val_mean_squared_error: 195.7875
Epoch 71/250
26/26 _____ 0s 7ms/step - loss: 207.0093 -
mean_absolute_error: 10.4519 - mean_squared_error: 206.6990 -
val_loss: 199.4959 - val_mean_absolute_error: 9.8384 -
val_mean_squared_error: 199.1850
Epoch 72/250
26/26 _____ 0s 7ms/step - loss: 213.9744 -
mean_absolute_error: 10.7788 - mean_squared_error: 213.6633 -
val_loss: 192.4788 - val_mean_absolute_error: 9.6919 -
val_mean_squared_error: 192.1671
Epoch 73/250
26/26 _____ 0s 8ms/step - loss: 183.1068 -
mean_absolute_error: 10.4433 - mean_squared_error: 182.7949 -
val_loss: 197.3482 - val_mean_absolute_error: 9.7927 -
val_mean_squared_error: 197.0358
Epoch 74/250
26/26 _____ 0s 7ms/step - loss: 203.7340 -
mean_absolute_error: 10.4998 - mean_squared_error: 203.4215 -
```



```
val_loss: 193.2769 - val_mean_absolute_error: 9.7967 -  
val_mean_squared_error: 192.9640  
Epoch 75/250  
26/26 ━━━━━━━━━━━ 0s 7ms/step - loss: 173.9892 -  
mean_absolute_error: 9.7234 - mean_squared_error: 173.6761 - val_loss:  
192.7126 - val_mean_absolute_error: 9.6827 - val_mean_squared_error:  
192.3991  
Epoch 76/250  
26/26 ━━━━━━━━━━━ 0s 11ms/step - loss: 166.4318 -  
mean_absolute_error: 9.6195 - mean_squared_error: 166.1181 - val_loss:  
195.2146 - val_mean_absolute_error: 9.7071 - val_mean_squared_error:  
194.9005  
Epoch 77/250  
26/26 ━━━━━━━━━━━ 0s 11ms/step - loss: 167.6431 -  
mean_absolute_error: 9.7513 - mean_squared_error: 167.3288 - val_loss:  
197.1027 - val_mean_absolute_error: 9.8481 - val_mean_squared_error:  
196.7880  
Epoch 78/250  
26/26 ━━━━━━━━━━━ 1s 13ms/step - loss: 190.0085 -  
mean_absolute_error: 10.3525 - mean_squared_error: 189.6936 -  
val_loss: 193.9883 - val_mean_absolute_error: 9.6975 -  
val_mean_squared_error: 193.6729  
Epoch 79/250  
26/26 ━━━━━━━━━━━ 0s 12ms/step - loss: 195.1637 -  
mean_absolute_error: 10.4491 - mean_squared_error: 194.8483 -  
val_loss: 196.9236 - val_mean_absolute_error: 9.8235 -  
val_mean_squared_error: 196.6077  
Epoch 80/250  
26/26 ━━━━━━━━━━━ 1s 11ms/step - loss: 162.8532 -  
mean_absolute_error: 9.3572 - mean_squared_error: 162.5371 - val_loss:  
198.4034 - val_mean_absolute_error: 9.9580 - val_mean_squared_error:  
198.0868  
Epoch 81/250  
26/26 ━━━━━━━━━━━ 0s 7ms/step - loss: 187.5797 -  
mean_absolute_error: 10.0736 - mean_squared_error: 187.2630 -  
val_loss: 193.2899 - val_mean_absolute_error: 9.6366 -  
val_mean_squared_error: 192.9726  
Epoch 82/250  
26/26 ━━━━━━━━━━━ 0s 8ms/step - loss: 181.5526 -  
mean_absolute_error: 10.0136 - mean_squared_error: 181.2352 -  
val_loss: 193.5690 - val_mean_absolute_error: 9.7739 -  
val_mean_squared_error: 193.2511  
  
plot_model(model_func, show_shapes=True, dpi=75)
```

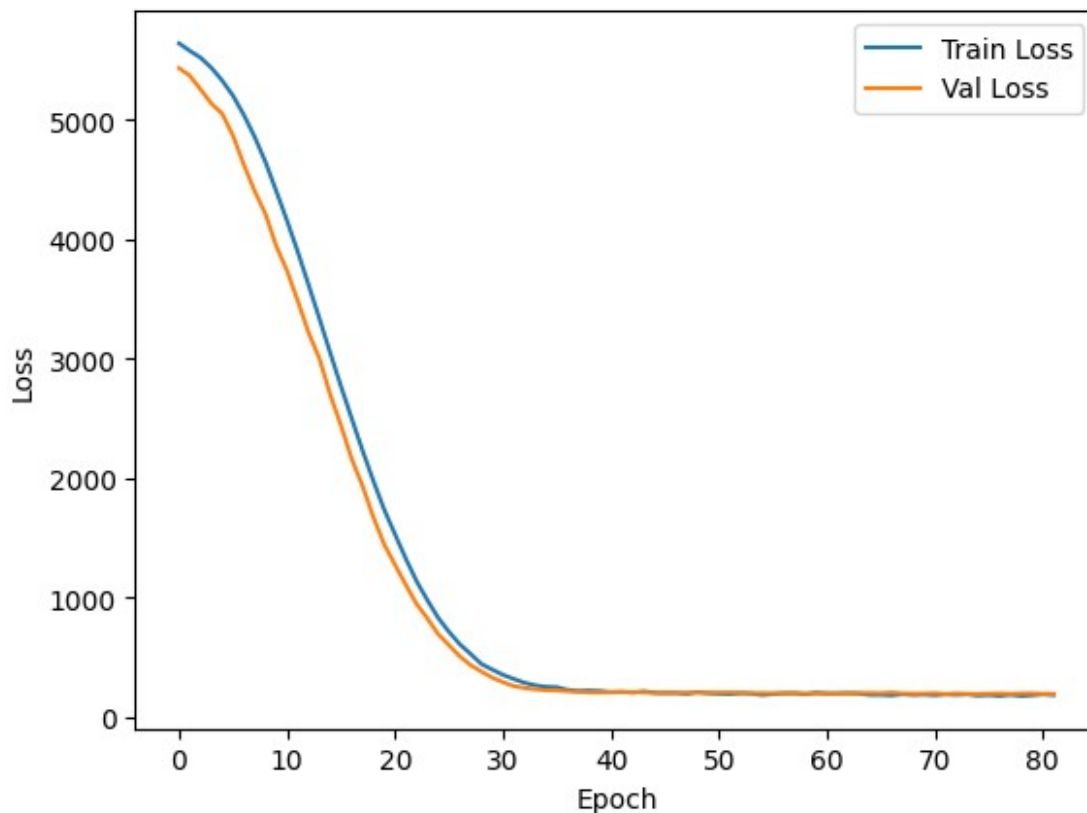


```
y_pred = model_func.predict(x_test_scaled)
```

4/4 ————— 0s 45ms/step

Evaluate

```
train_loss = history.history['loss']
val_loss = history.history['val_loss']
plt.plot(train_loss, label="Train Loss")
plt.plot(val_loss, label='Val Loss')
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



Based on the plot, the loss curve for both training and validation in modified functional API model shows a consistent decrease throughout epochs, which means the model is effectively learning from the training data while generating well with the validation dataset. The small gap between the two curves suggests that the model is not overfitting, and it can be concluded that it has converged by epoch 60.

```
R2 = r2_score(y_test, y_pred)
MAE = mean_absolute_error(y_test, y_pred)
MSE = mean_squared_error(y_test, y_pred)
```

```

print("R2 Score=", R2)
print("Mean Absolute Error=", MAE)
print("Mean Squared Error=", MSE)

R2 Score= 0.30773368610369567
Mean Absolute Error= 11.077980034795301
Mean Squared Error= 245.13486270487005

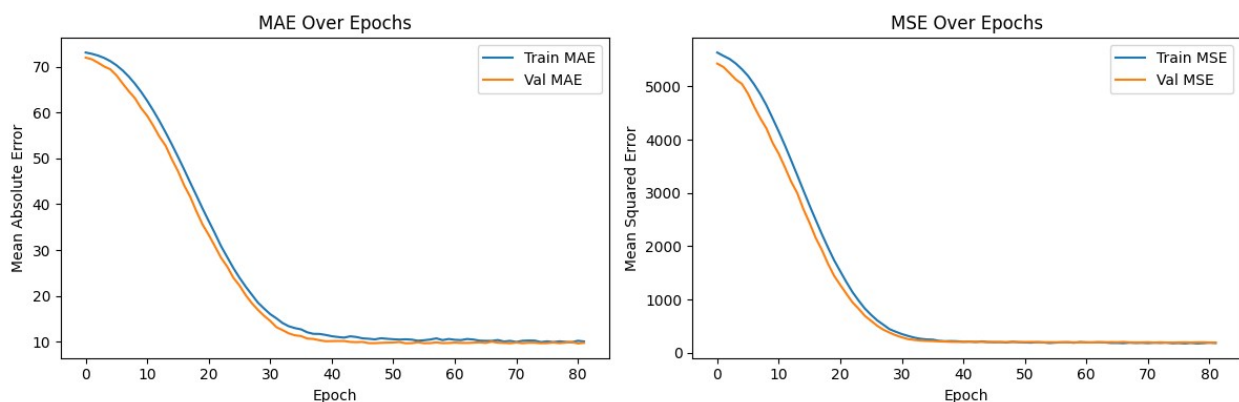
# Plot Mean Absolute Error (MAE)
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.plot(history.history['mean_absolute_error'], label='Train MAE')
plt.plot(history.history['val_mean_absolute_error'], label='Val MAE')
plt.xlabel('Epoch')
plt.ylabel('Mean Absolute Error')
plt.title('MAE Over Epochs')
plt.legend()

# Plot Mean Squared Error (MSE)
plt.subplot(1, 2, 2)
plt.plot(history.history['mean_squared_error'], label='Train MSE')
plt.plot(history.history['val_mean_squared_error'], label='Val MSE')
plt.xlabel('Epoch')
plt.ylabel('Mean Squared Error')
plt.title('MSE Over Epochs')
plt.legend()

plt.tight_layout()
plt.show()

```



The statistical results for modified functional API model shows a low values of R^2 (0.3), and high values of MAE (11.07) and MSE (245.13). This means that the model's predictivity is limited because its only 30% of the variance in the data.

But based on the curve plots, both MSE and MAE decrease consistently throughout epochs, which means the model learn effectively without overfitting.

Comparison Both Modified Model

Both modified models shows that the loss curve decrease consistently, we can indicate that the models are learning effectively from the dataset and showing no overfitting. However, based on the statistical results such as values of R^2 , MAE, and MSE, suggest that the model are not yet performing optimal predict for productivity score, as both values for R^2 is low, and high values of MAE and MSE.

Additionally, when we compare both modified model, the functional API model show better performance than the sequential model.

Conclusion

Based on the goal to estimate the productivity score for each team in a garment factory, I went through a comprehensive process involving data cleaning, visualization, preprocessing, model building, and evaluation. At first, both the baseline models (sequential and functional API) showed they learned fast because the loss curve were decreased in small epochs, but there is no indicate overfitting. While this was quite good, the results weren't good based on the statistical results with low R^2 and high MAE and MSE.

After modifying the models by adding batch normalization, dropout, change ReLU into ELU activation, and apply L2 regulation, I notice a better performance. The loss curve for both models are decrease consistently without any overfitting, which means the models were learning effectively. They even converged around epochs 50/60.

However, all statistical results still shows low values of R^2 and high values for both MAE and MSE, which means need to keep improving. This also might because of the dataset has its own challenges, such as quality issues or insufficient diversity in samples.

Presentation Video Link :

<https://drive.google.com/file/d/1RJrgK1A4kKmcwhBMtvc1UYC9NBNQFPIV/view?usp=sharing>